

# CS 2420: Binary Heap

---

Dr. Tsung-Wei (TW) Huang

Department of Electrical and Computer Engineering

University of Utah, Salt Lake City, UT



# Outline

---

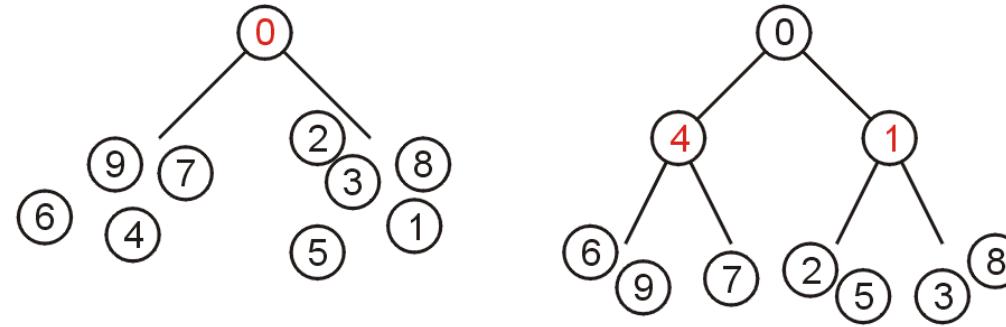
- In this topic, we will:
  - Define a binary min-heap
  - Look at some examples
  - Operations on heaps:
    - Top
    - Pop
    - Push
  - An array representation of heaps
  - Define a binary max-heap
  - Using binary heaps as priority queues

# Definition

---

A non-empty binary tree is a min-heap if

- The key associated with the root is less than or equal to the keys associated with either of the sub-trees (if any)
- Both of the sub-trees (if any) are also binary min-heaps



From this definition:

- A single node is a min-heap
- All keys in either sub-tree are greater than the root key

# Definition

---

Important:

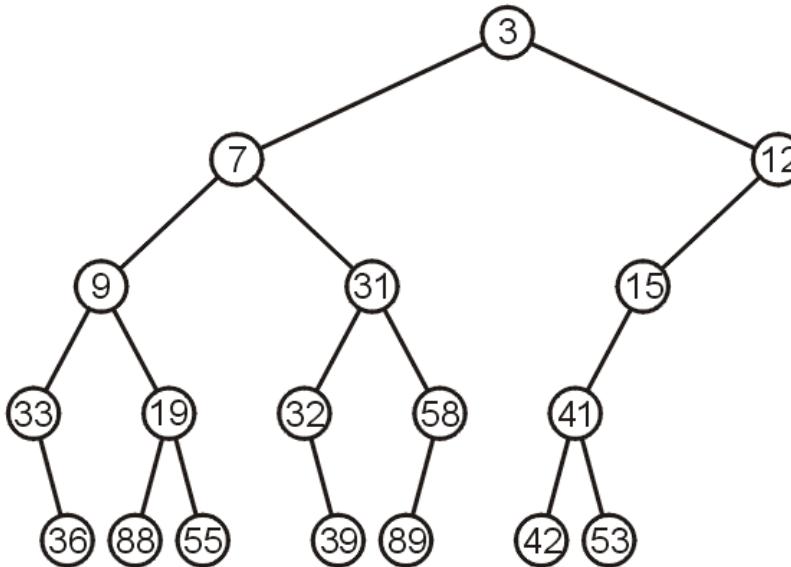
**THERE IS NO OTHER RELATIONSHIP BETWEEN  
THE ELEMENTS IN THE TWO SUBTREES**

Failing to understand this is the greatest mistake a student makes

# Example

---

This is a binary min-heap:

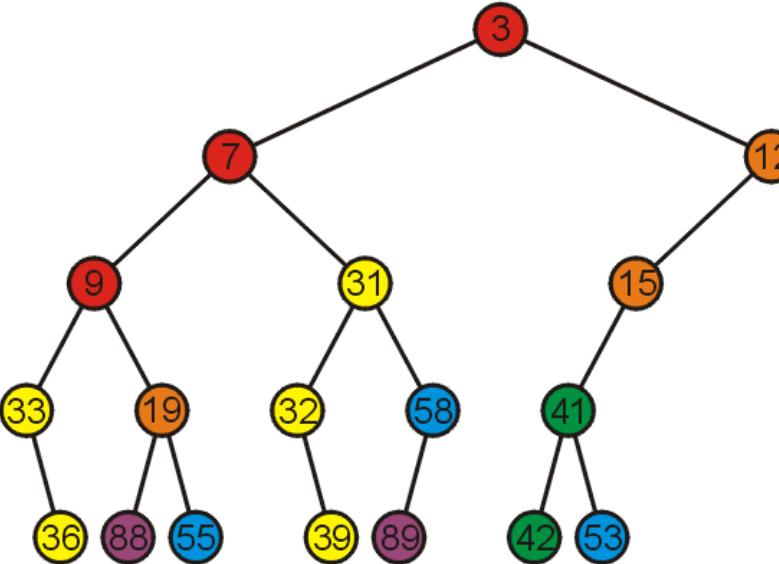


# Example

---

Adding colour, we observe

- The left subtree has the smallest (7) and the largest (89) objects
- No relationship between items with similar priority



# Operations

---

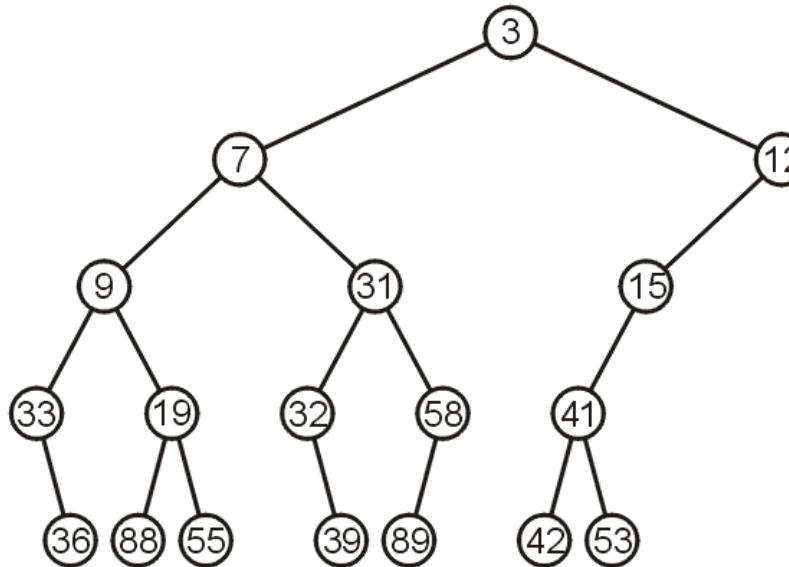
We will consider three operations:

- Top
- Pop
- Push

# Example

---

We can find the top object in  $\Theta(1)$  time: 3



# Pop

---

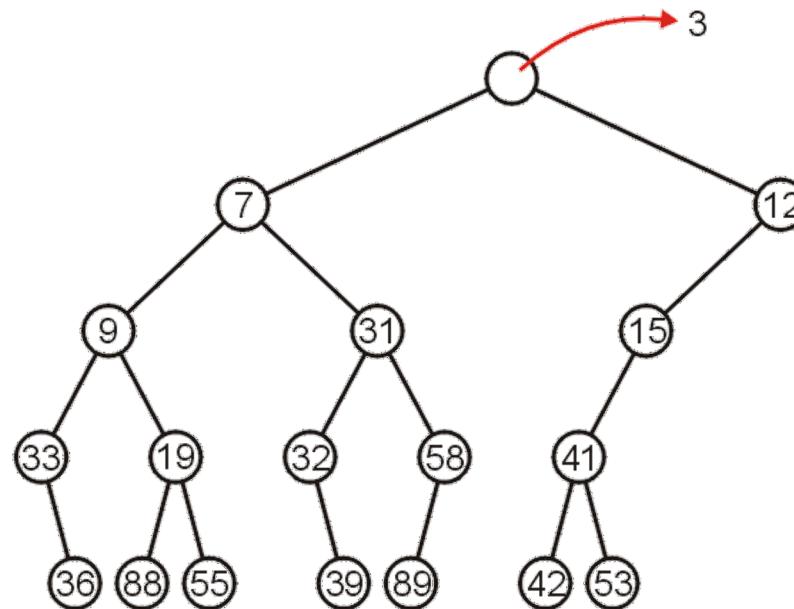
To remove the minimum object:

- Promote the node of the sub-tree which has the least value
- Recurs down the sub-tree from which we promoted the least value

# Pop

---

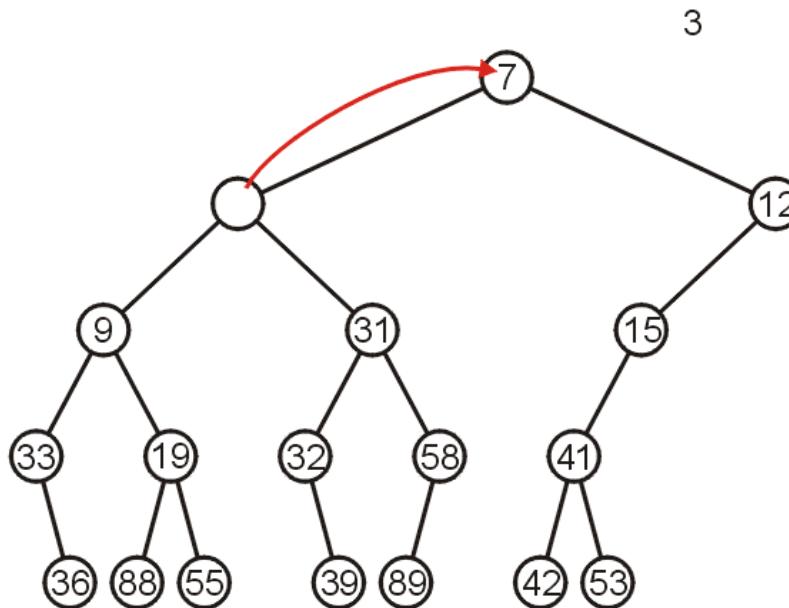
Using our example, we remove 3:



# Pop

---

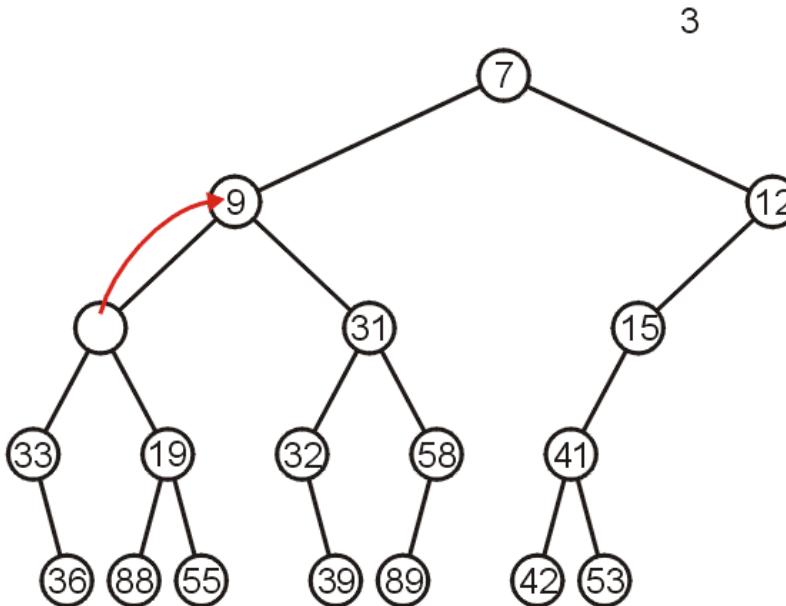
We promote 7 (the minimum of 7 and 12) to the root:



# Pop

---

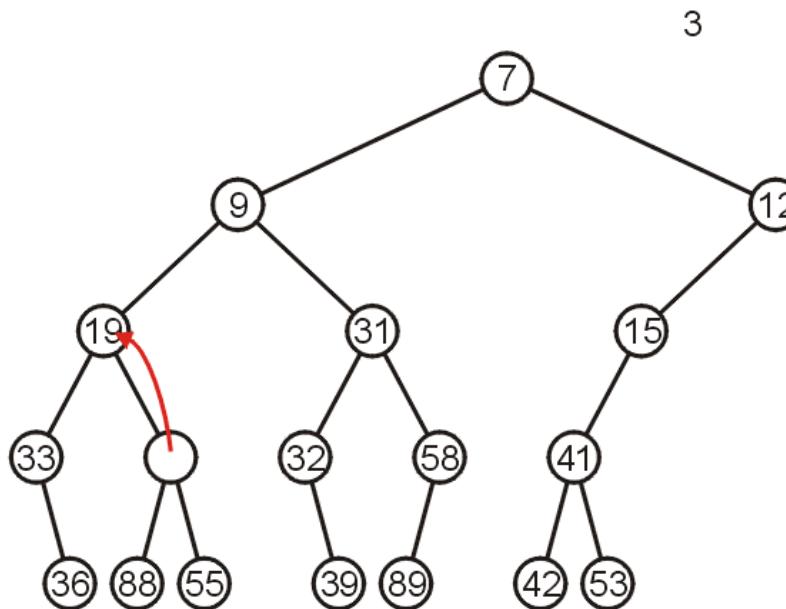
In the left sub-tree, we promote 9:



# Pop

---

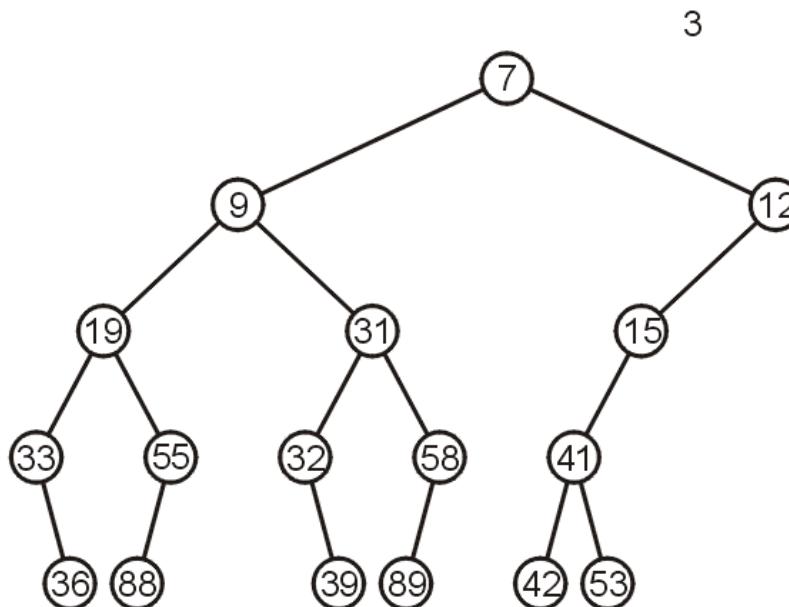
Recursively, we promote 19:



# Pop

---

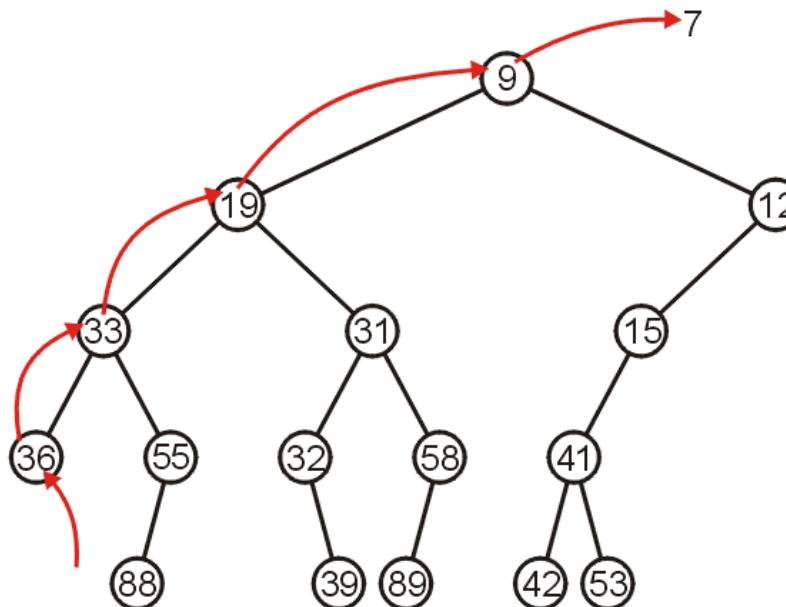
Finally, 55 is a leaf node, so we promote it and delete the leaf



# Pop

---

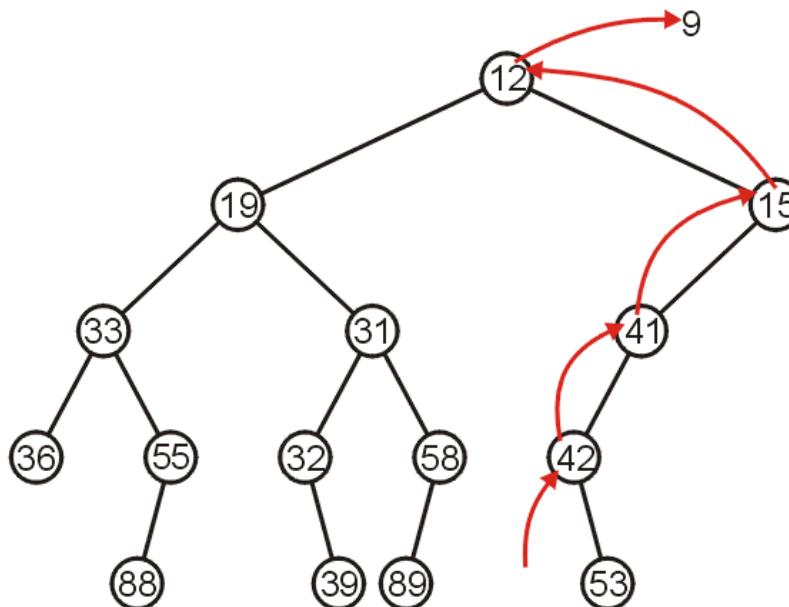
Repeating this operation again, we can remove 7:



# Pop

---

If we remove 9, we must now promote from the right sub-tree:



# Push

---

Inserting into a heap may be done either:

- At a leaf (move it up if it is smaller than the parent)
- At the root (insert the larger object into one of the subtrees)

We will use the first approach with binary heaps

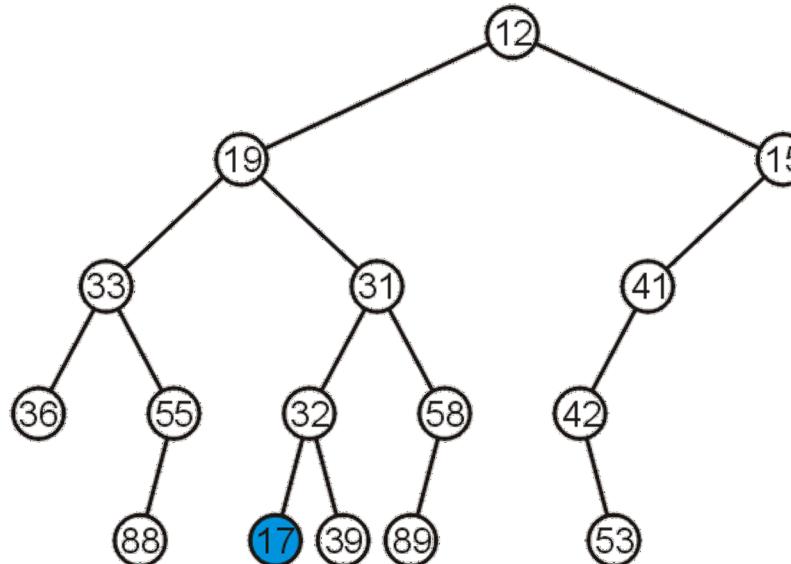
- Other heaps use the second

# Push

---

Inserting 17 into the last heap

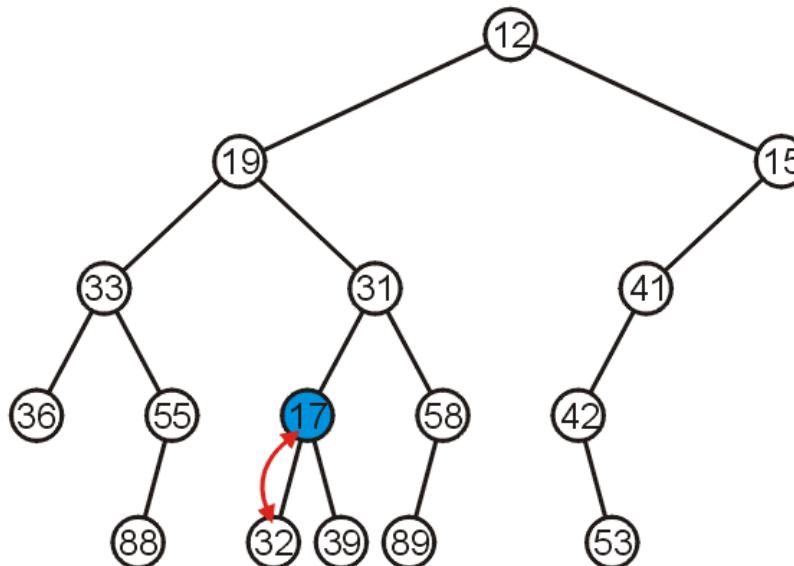
- Select an arbitrary node to insert a new leaf node:



# Push

---

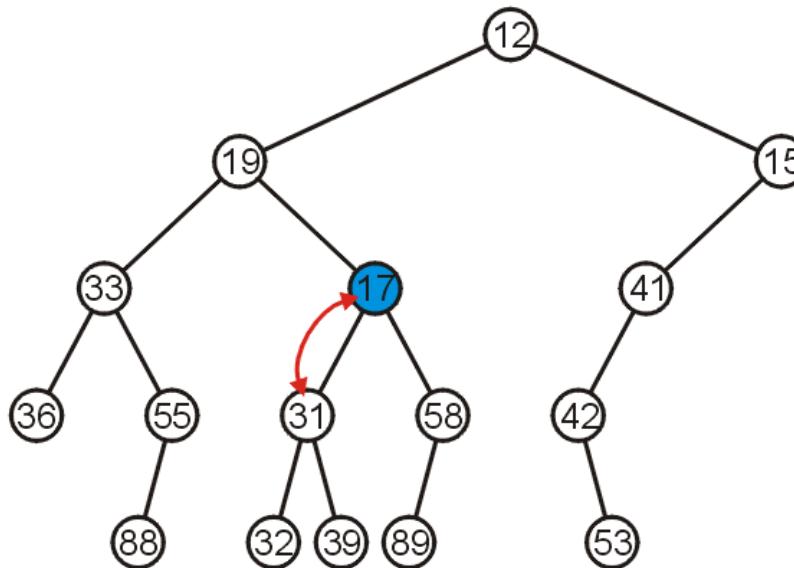
The node 17 is less than the node 32, so we swap them



# Push

---

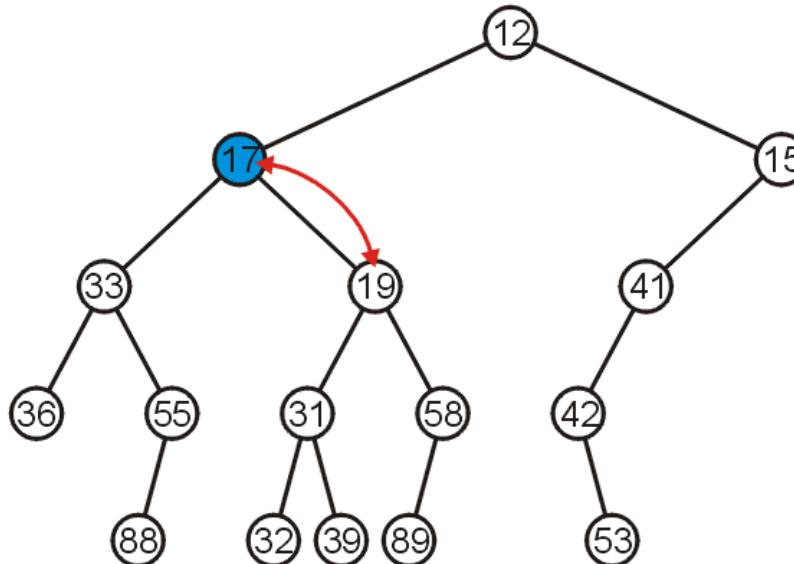
The node 17 is less than the node 31; swap them



# Push

---

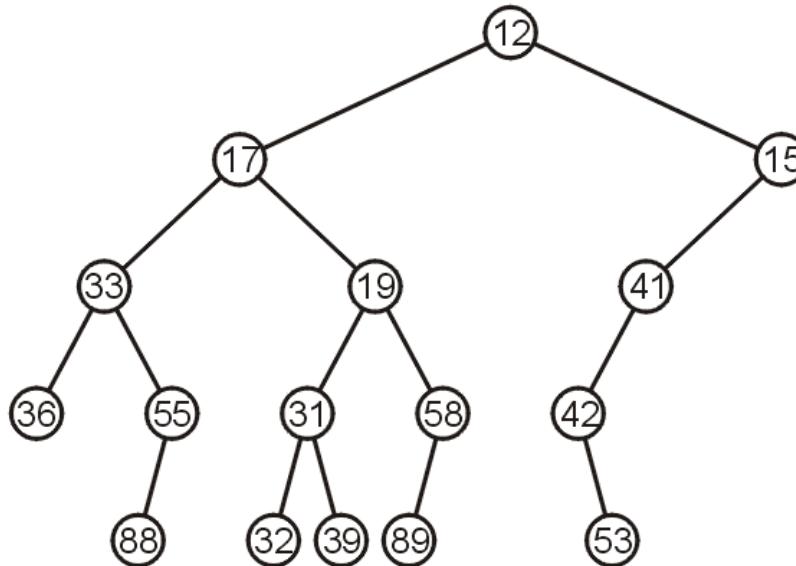
The node 17 is less than the node 19; swap them



# Push

---

The node 17 is greater than 12 so we are finished



# Push

---

Observation: both the left and right subtrees of 19 were greater than 19, thus we are guaranteed that we don't have to send the new node down

This process is called *percolation*, that is, the lighter (smaller) objects move up from the bottom of the min-heap

# Implementations

---

With binary search trees, we introduced the concept of *balance*

From this, we looked at:

- AVL Trees
- B-Trees
- Red-black Trees (not course material)

How can we determine where to insert so as to keep balance?

# Implementations

---

There are multiple means of keeping balance with binary heaps:

- Complete binary trees
- Leftist heaps
- Skew heaps

We will look at using complete binary trees

- It has optimal memory characteristics but sub-optimal run-time characteristics

# Complete Trees

---

By using complete binary trees, we will be able to maintain, with minimal effort, the complete tree structure

We have already seen

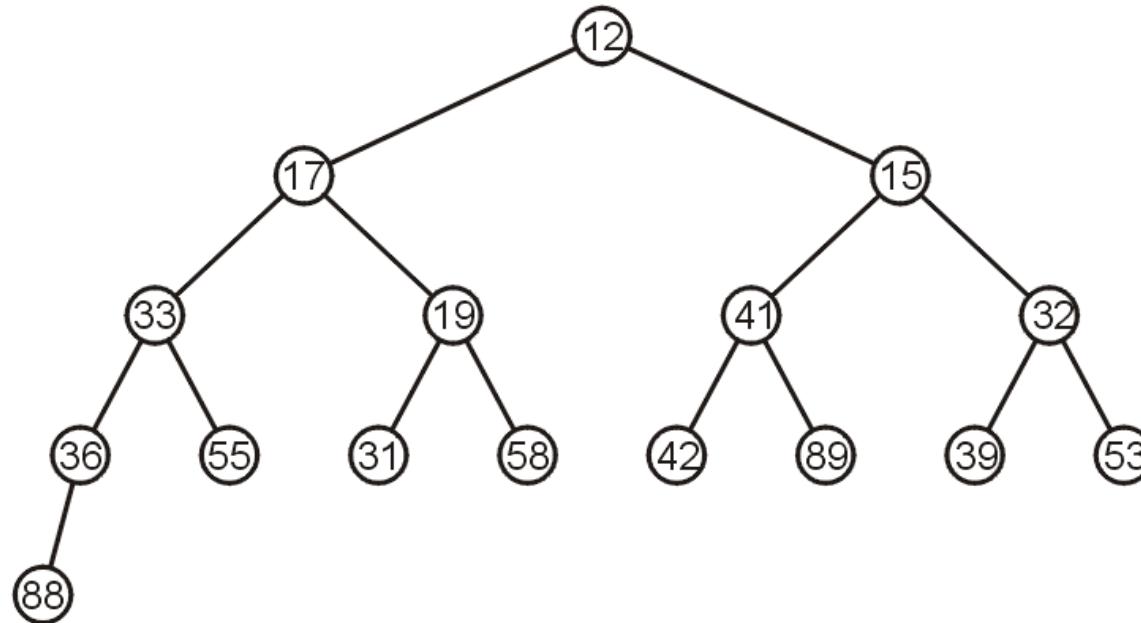
- It is easy to store a complete tree as an array

If we can store a heap of size  $n$  as an array of size  $\Theta(n)$ , this would be great!

# Complete Trees

---

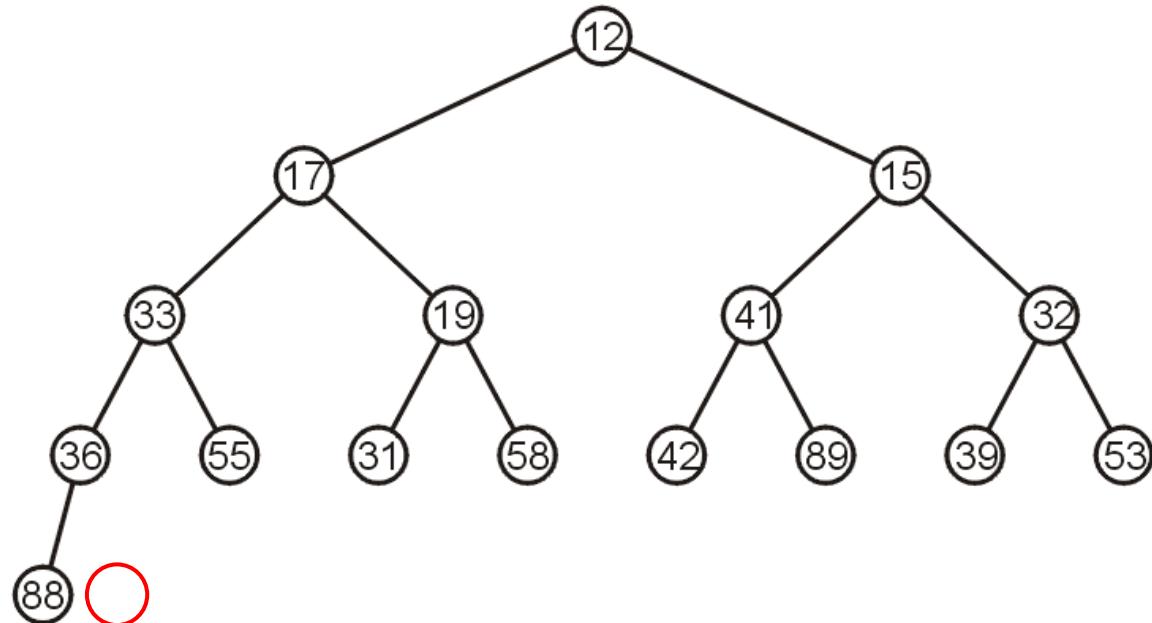
For example, the previous heap may be represented as the following (non-unique!) complete tree:



# Complete Trees: Push

---

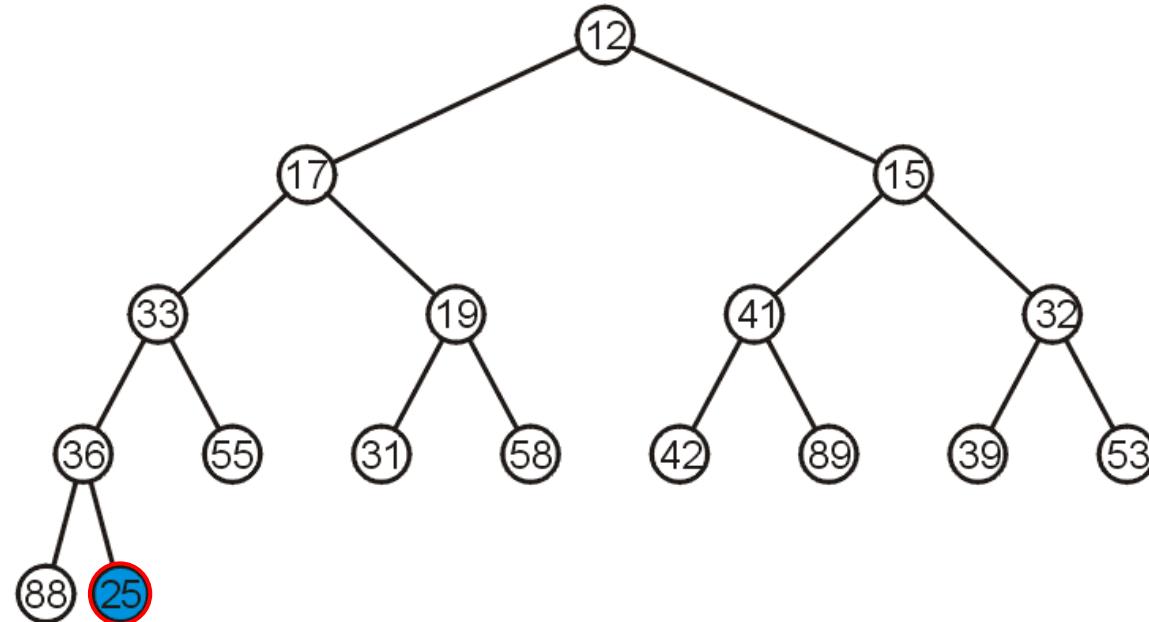
If we insert into a complete tree, we need only place the new node as a leaf node in the appropriate location and percolate up



# Complete Trees: Push

---

For example, push 25:

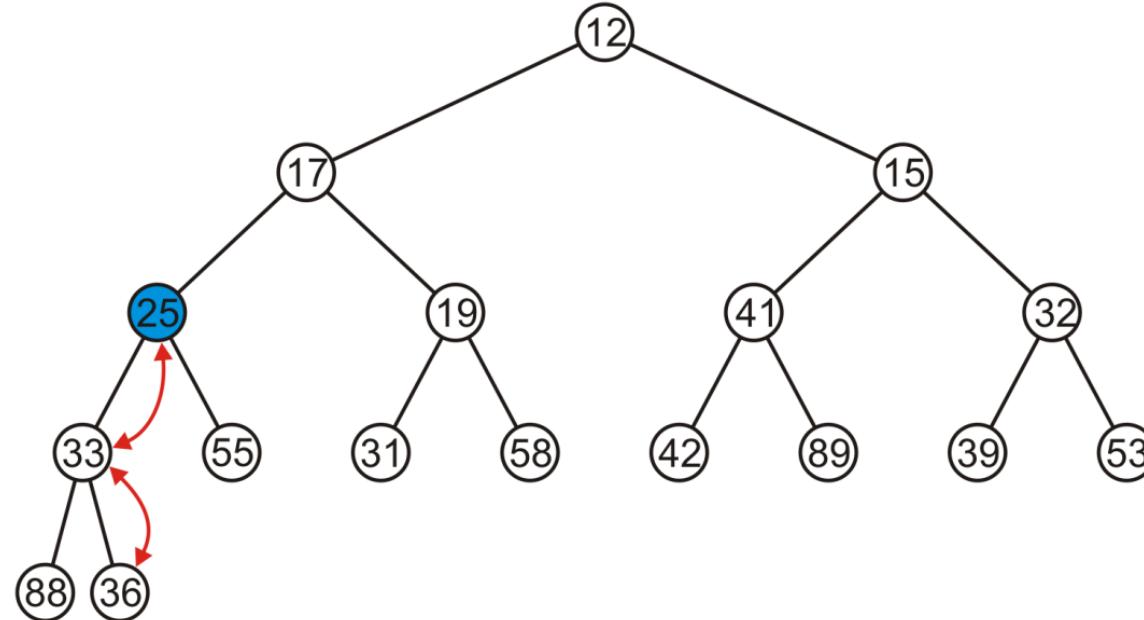


# Complete Trees: Push

---

We have to percolate 25 up into its appropriate location

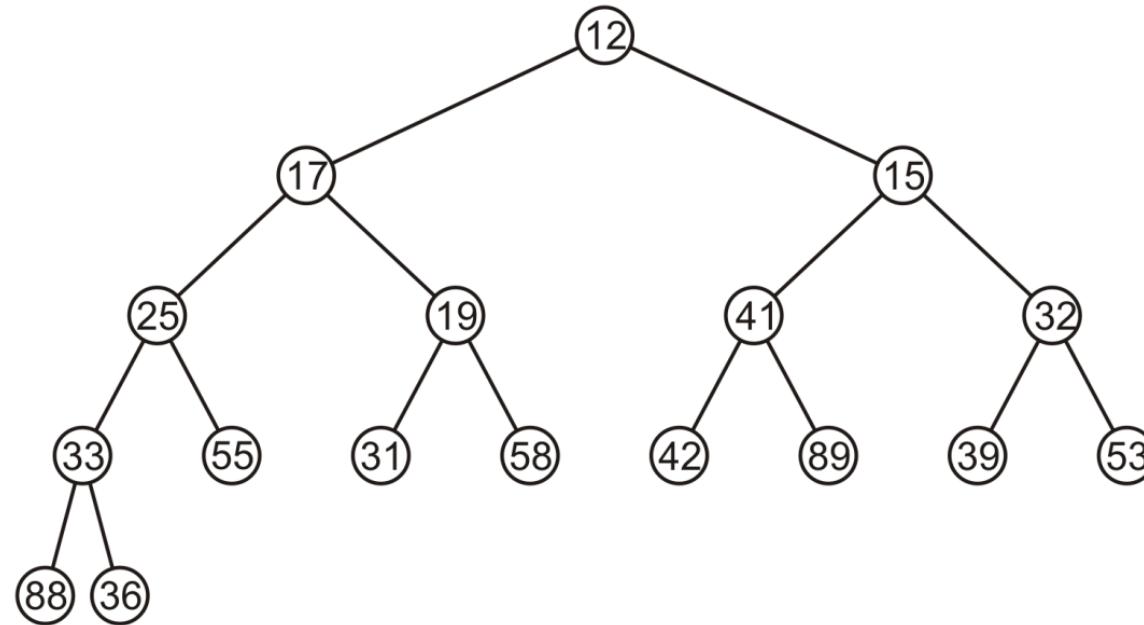
- The resulting heap is still a complete tree



# Complete Trees: Pop

---

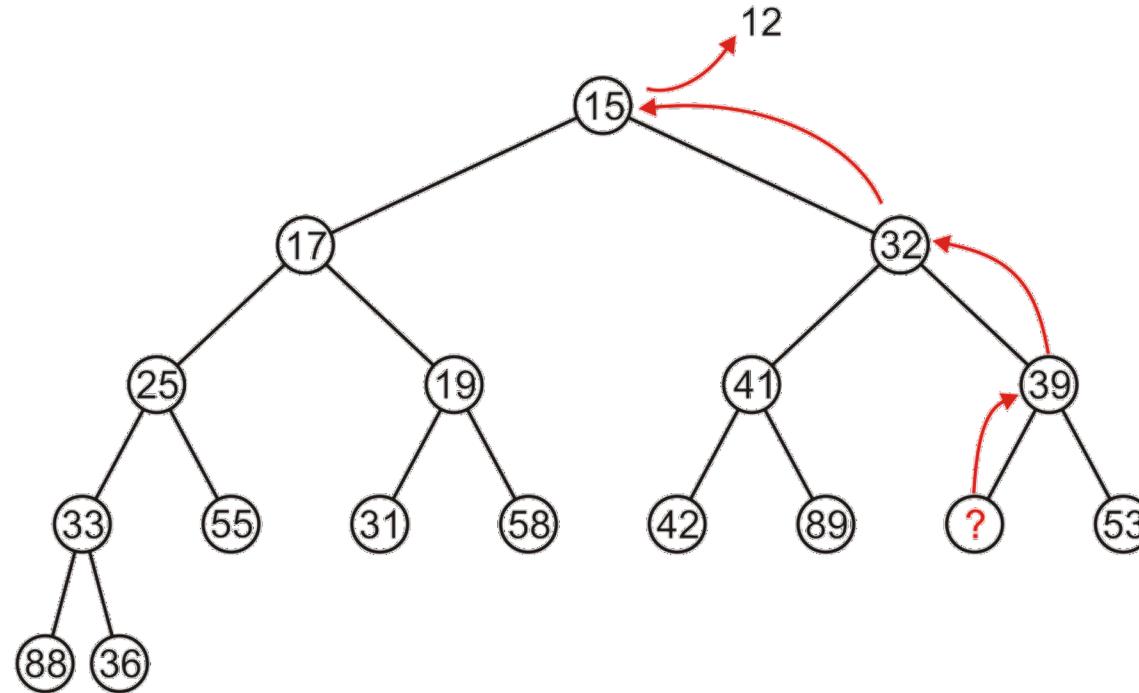
Suppose we want to pop the top entry: 12



# Complete Trees: Pop

---

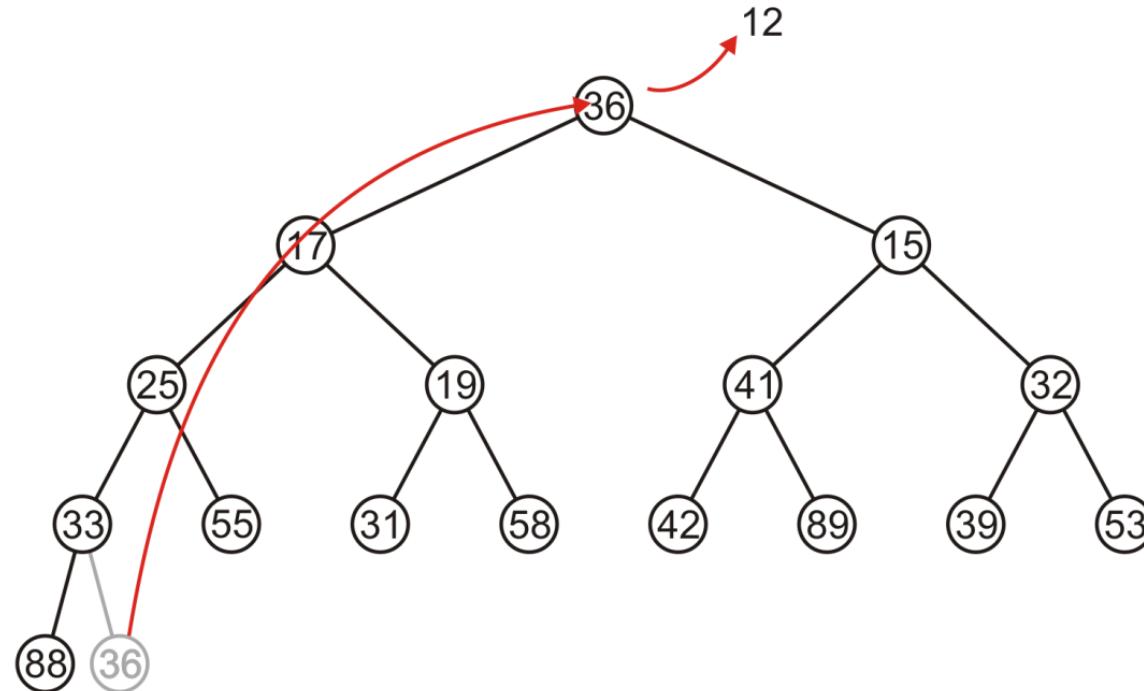
Percolating up creates a hole leading to a non-complete tree



# Complete Trees: Pop

---

Alternatively, copy the last entry in the heap to the root

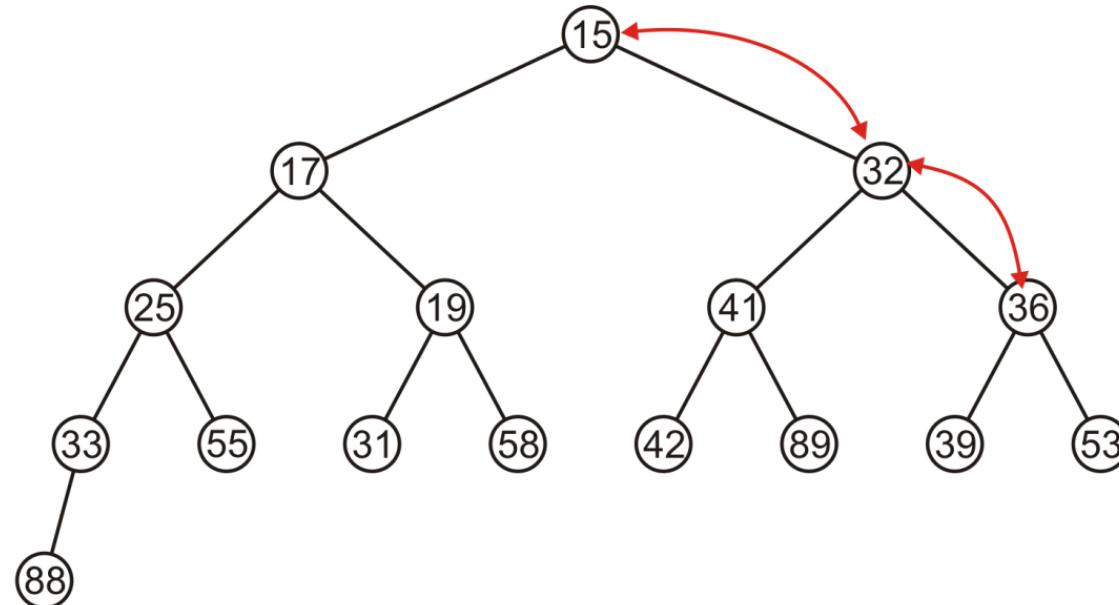


# Complete Trees: Pop

---

Now, percolate 36 down swapping it with the smallest of its children

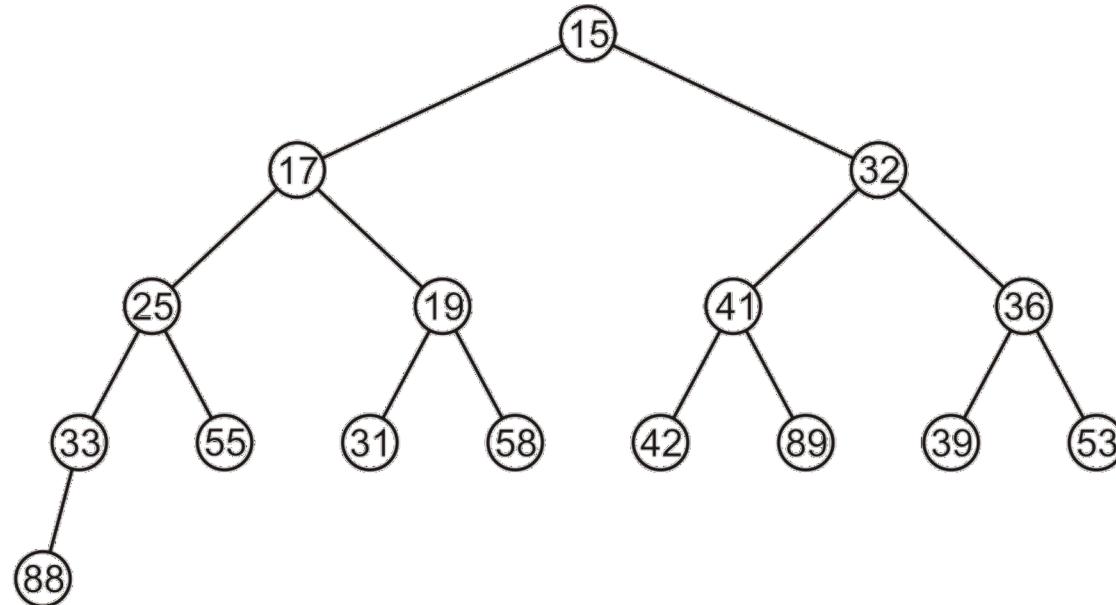
- We halt when both children are larger<sup>12</sup>



# Complete Trees: Pop

---

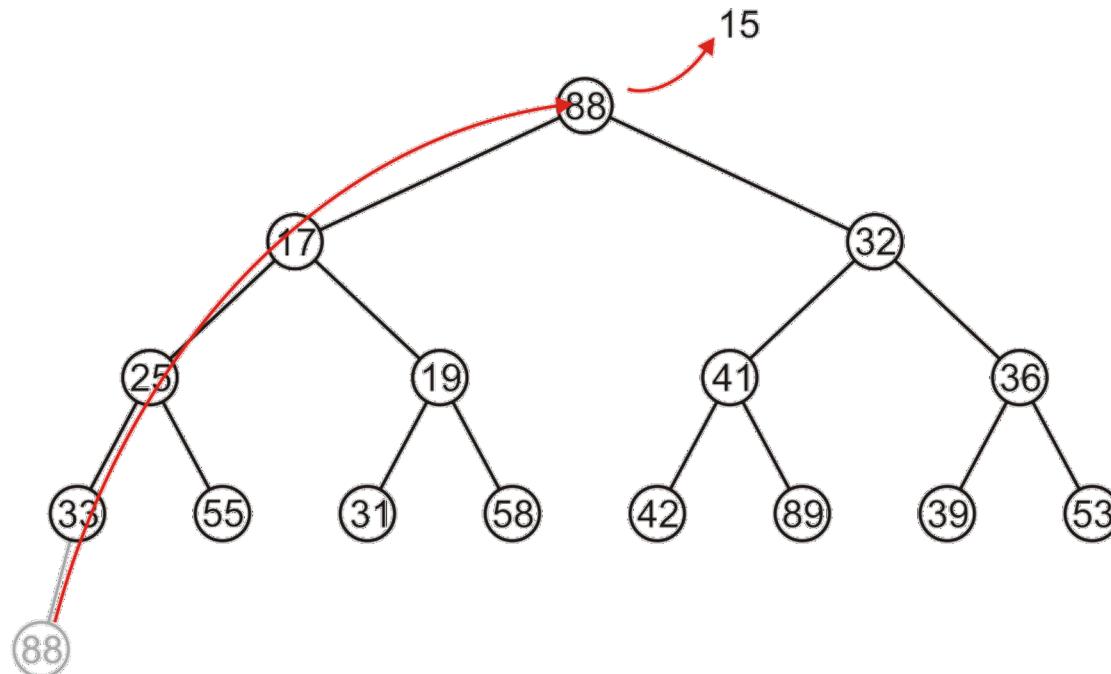
The resulting tree is now still a complete tree:



# Complete Trees: Pop

---

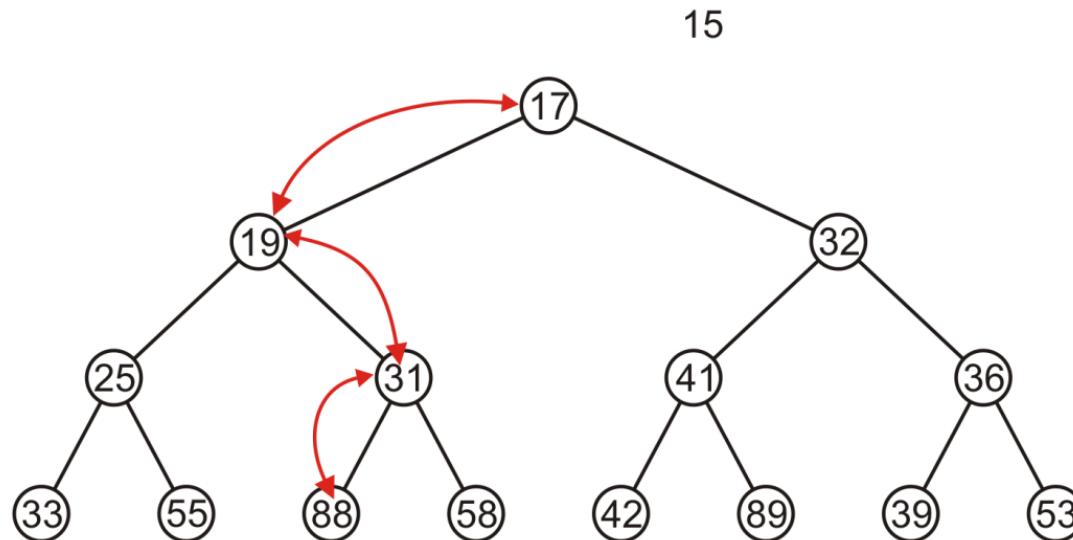
Again, popping 15, copy up the last entry: 88



# Complete Trees: Pop

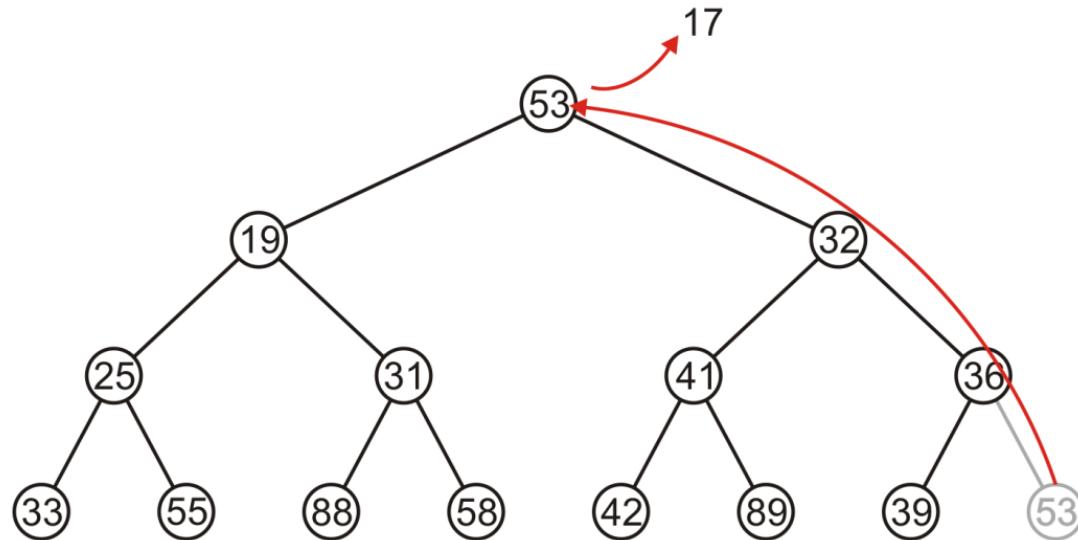
---

This time, it gets percolated down to the point where it has no children



# Complete Trees: Pop

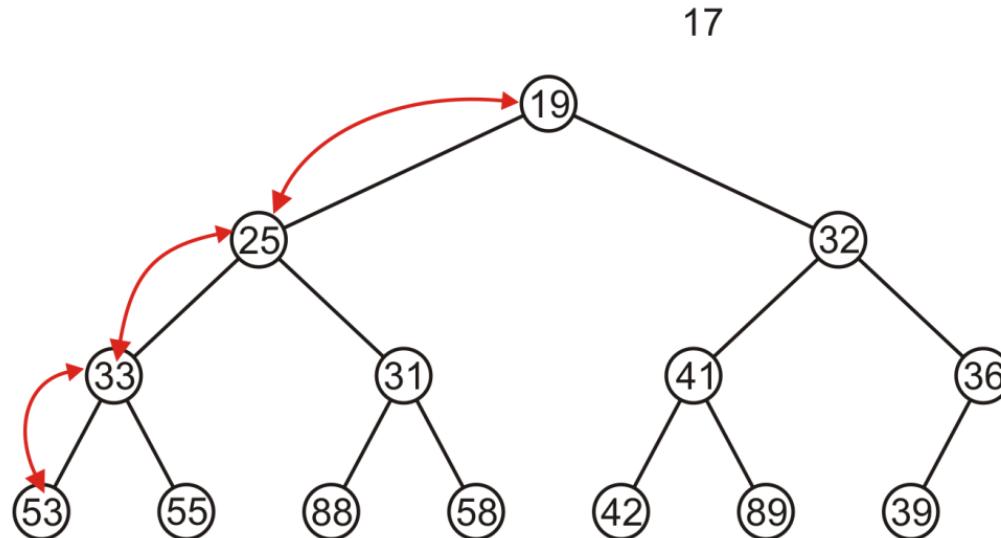
In popping 17, 53 is moved to the top



# Complete Trees: Pop

---

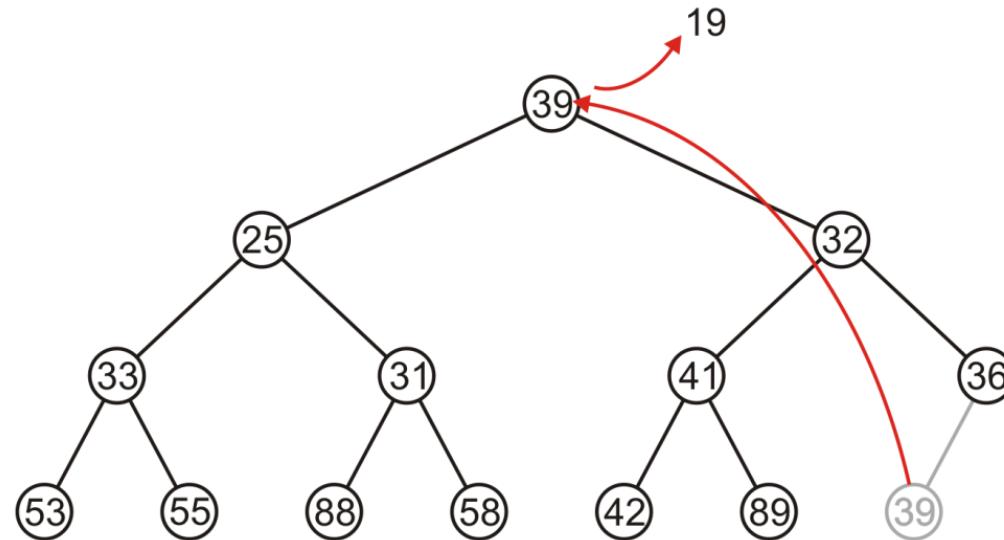
And percolated down, again to the deepest level



# Complete Trees: Pop

---

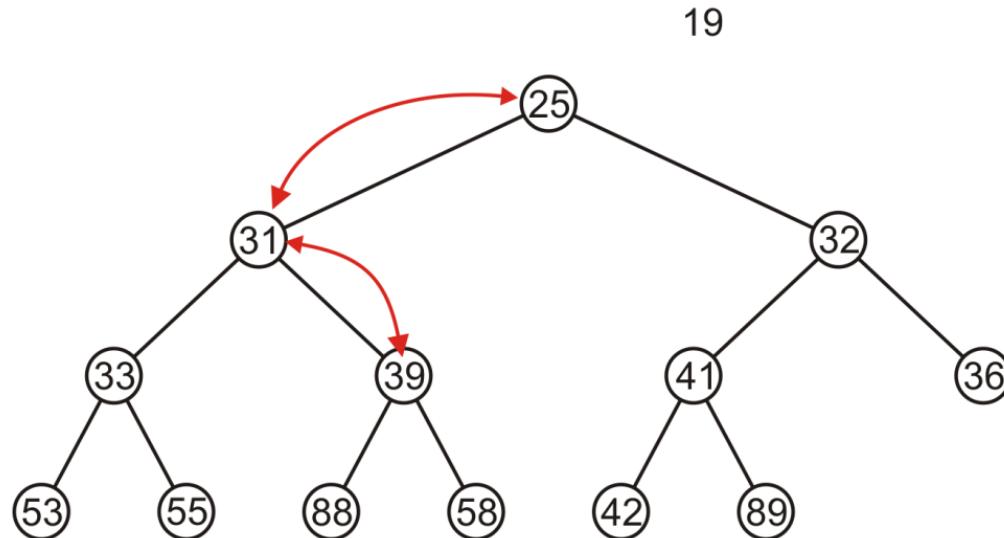
Popping 19 copies up 39



# Complete Trees: Pop

---

Which is then percolated down to the second deepest level



# Complete Tree

---

Therefore, we can maintain the complete-tree shape of a heap

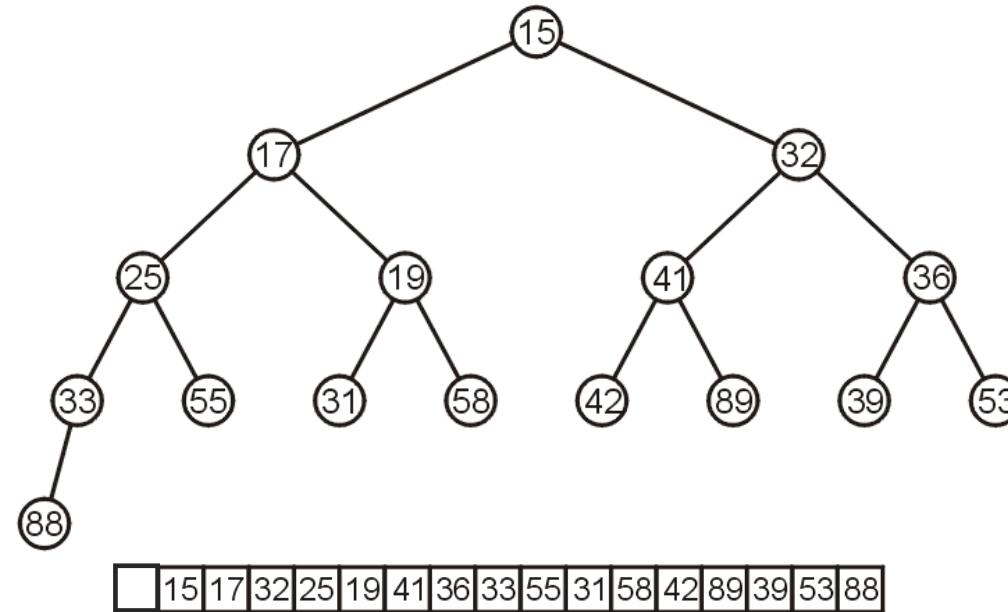
We may store a complete tree using an array:

- A complete tree is filled in breadth-first traversal order
- The array is filled using breadth-first traversal

# Array Implementation

---

For the heap

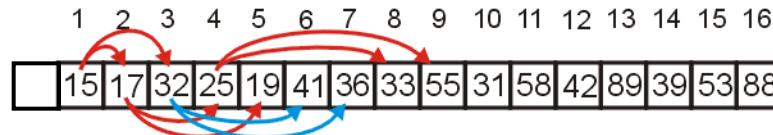


a breadth-first traversal yields:

# Array Implementation

---

Recall that if we associate an index—starting at 1—with each entry in the breadth-first traversal, we get:



Given the entry at index  $k$ , it follows that:

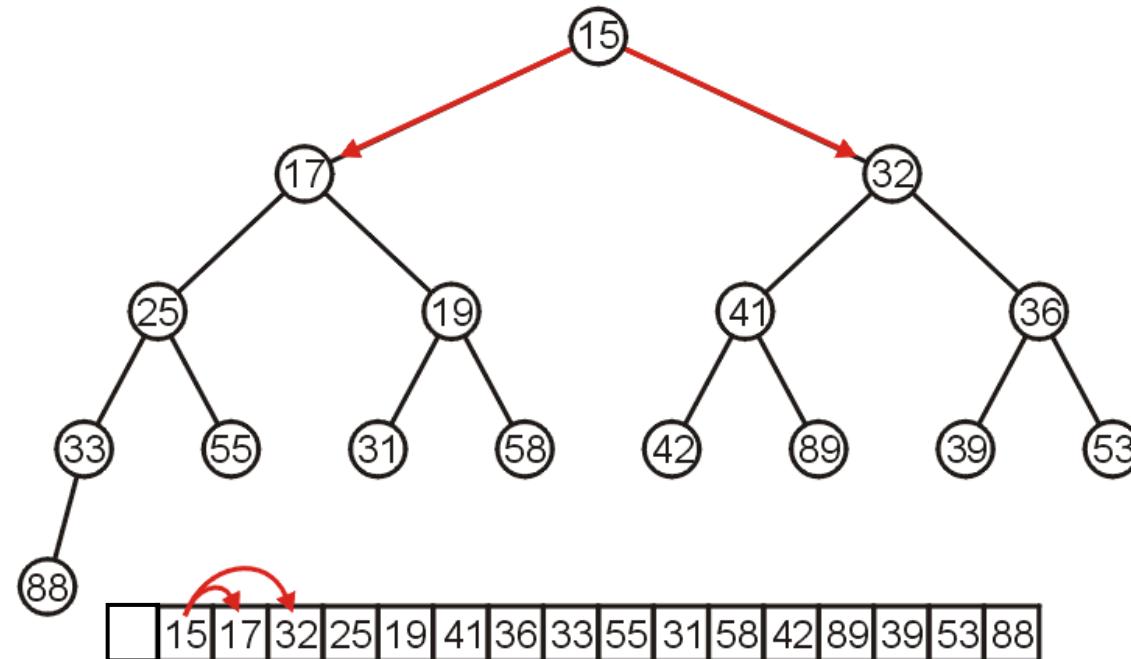
- The parent of node is at  $k/2$                            $\text{parent} = k \gg 1;$
- the children are at  $2k$  and  $2k + 1$                    $\text{left\_child} = k \ll 1;$   
     $\text{right\_child} = \text{left\_child} | 1;$

Cost (trivial): start array at position 1 instead of position 0

# Array Implementation

---

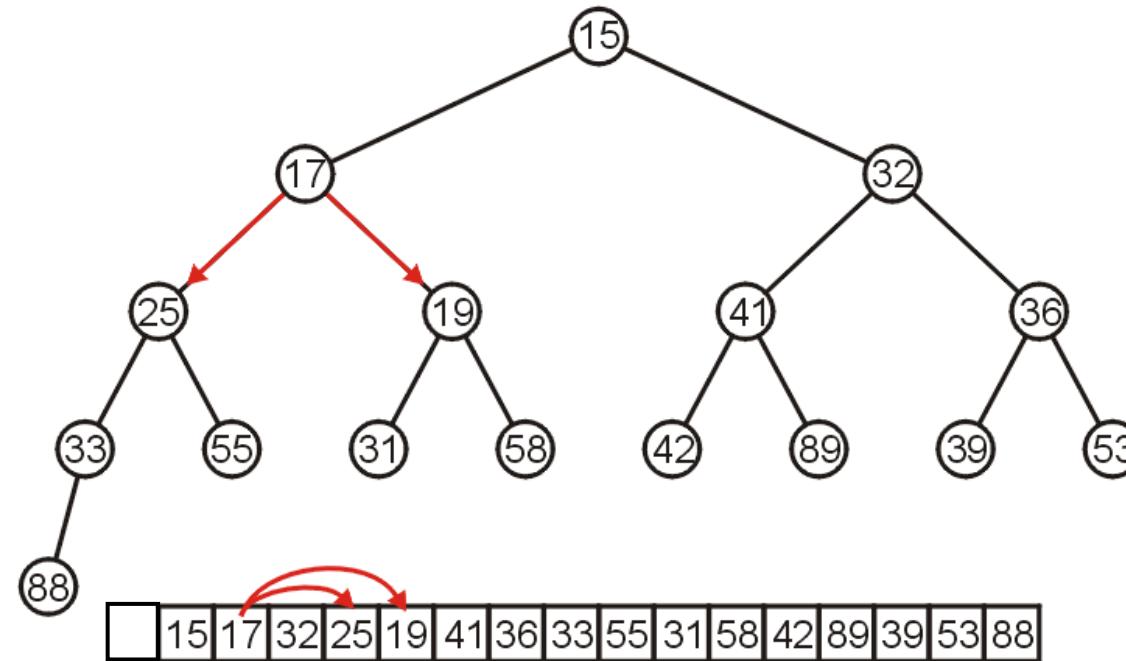
The children of 15 are 17 and 32:



# Array Implementation

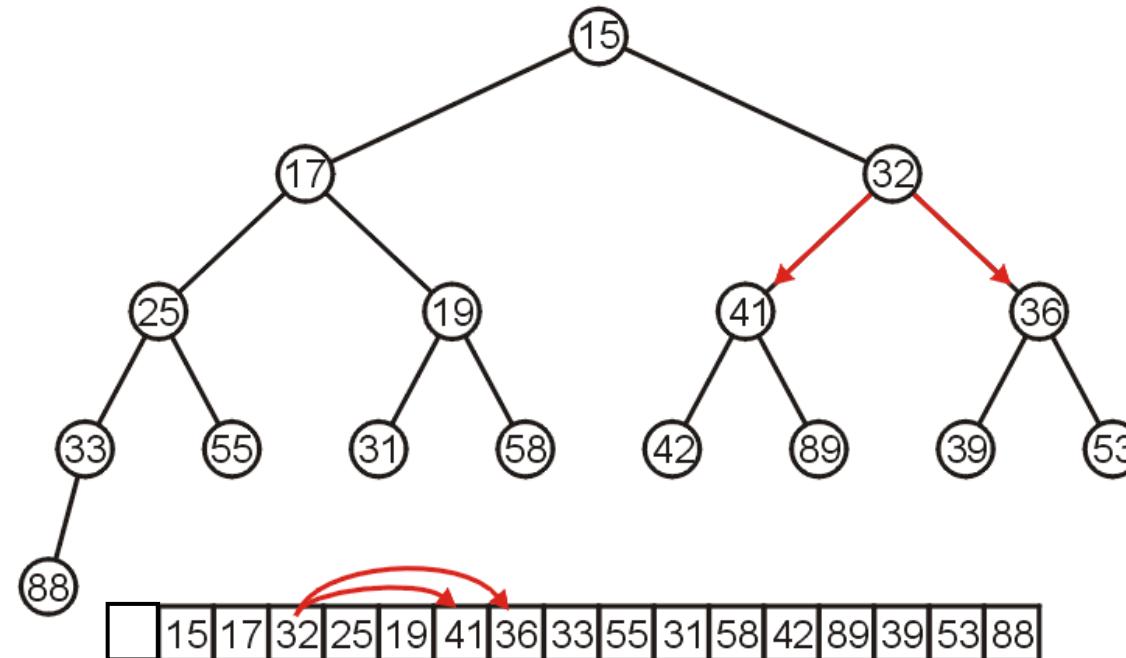
---

The children of 17 are 25 and 19:



# Array Implementation

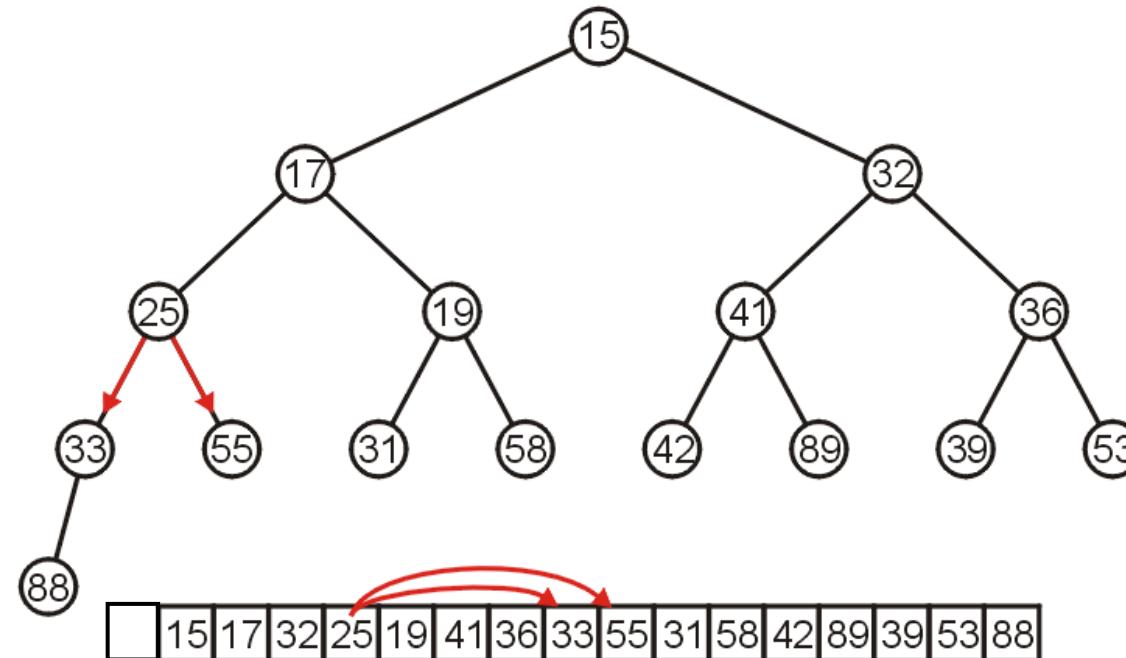
The children of 32 are 41 and 36:



# Array Implementation

---

The children of 25 are 33 and 55:



# Array Implementation

---

If the heap-as-array has **count** entries, then the next empty node in the corresponding complete tree is at location **posn = count + 1**

We compare the item at location **posn** with the item at **posn/2**

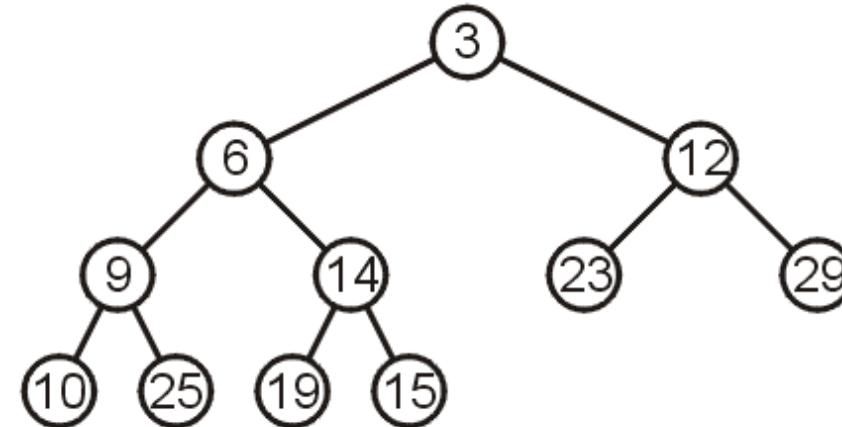
If they are out of order

- Swap them, set **posn /= 2** and repeat

# Array Implementation

---

Consider the following heap, both as a tree and in its array representation

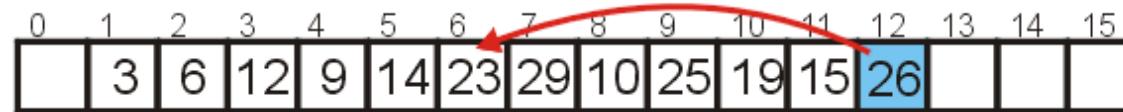
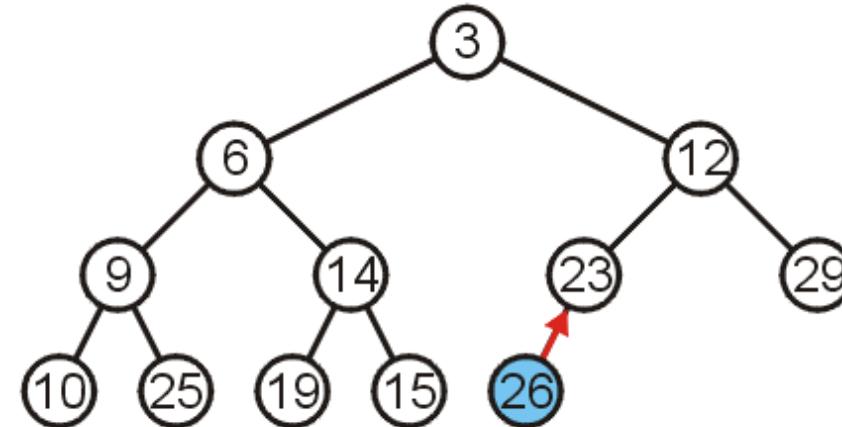


0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	3	6	12	9	14	23	29	10	25	19	15				

# Array Implementation: Push

---

Inserting 26 requires no changes



# Array Implementation: Push

---

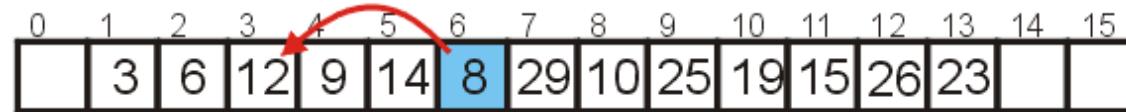
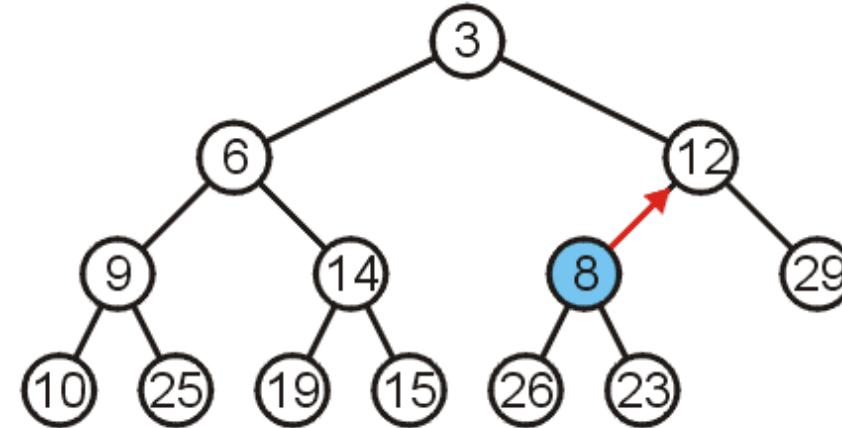
Inserting 8 requires a few percolations:

- Swap 8 and 23

# Array Implementation: Push

---

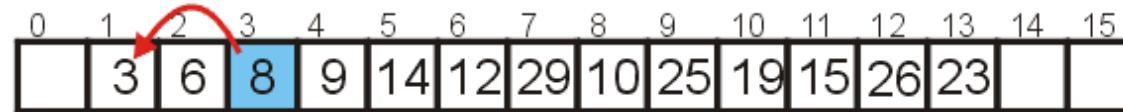
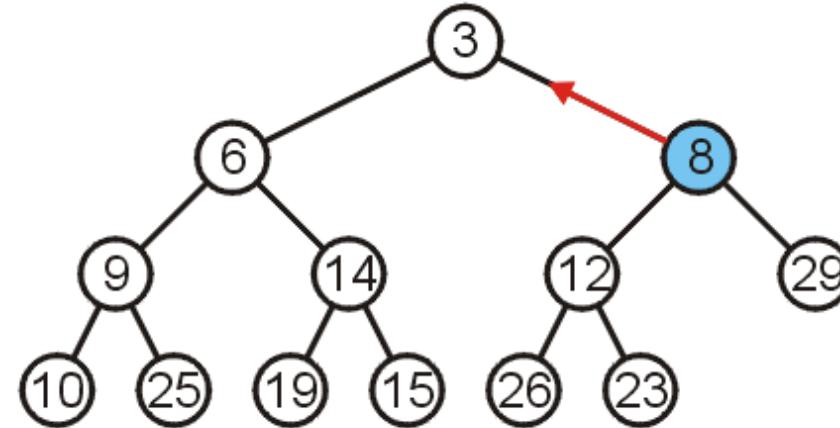
Swap 8 and 12



# Array Implementation: Push

---

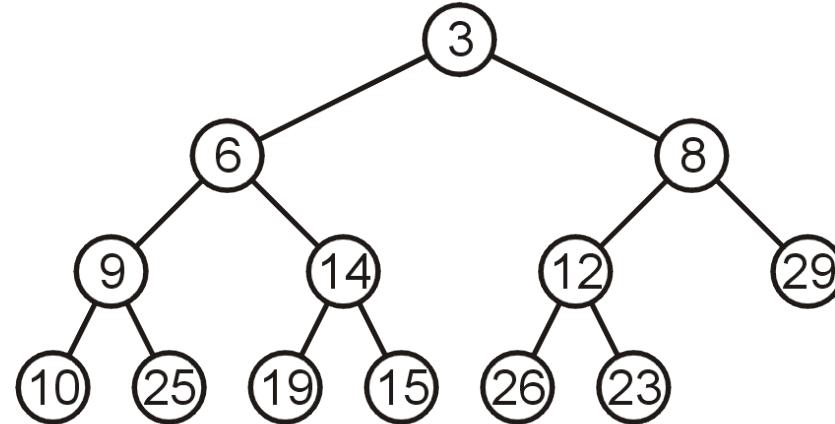
At this point, it is greater than its parent, so we are finished



# Array Implementation: Pop

---

As before, popping the top has us copy the last entry to the top



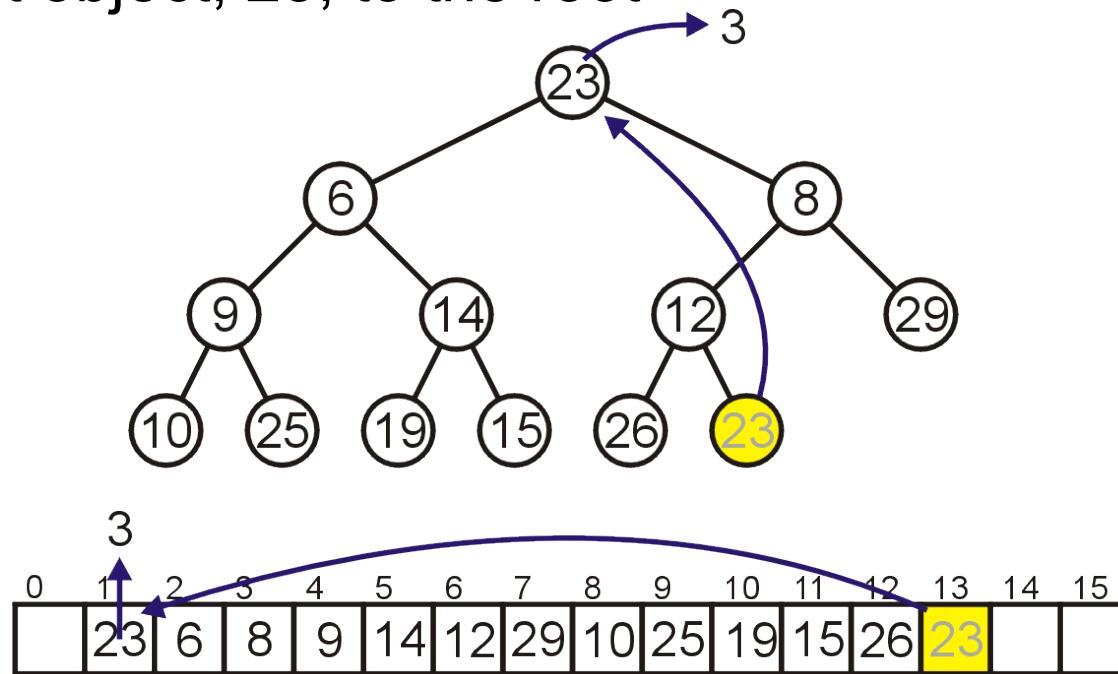
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	3	6	8	9	14	12	29	10	25	19	15	26	23		

# Array Implementation: Pop

---

Instead, consider this strategy:

- Copy the last object, 23, to the root



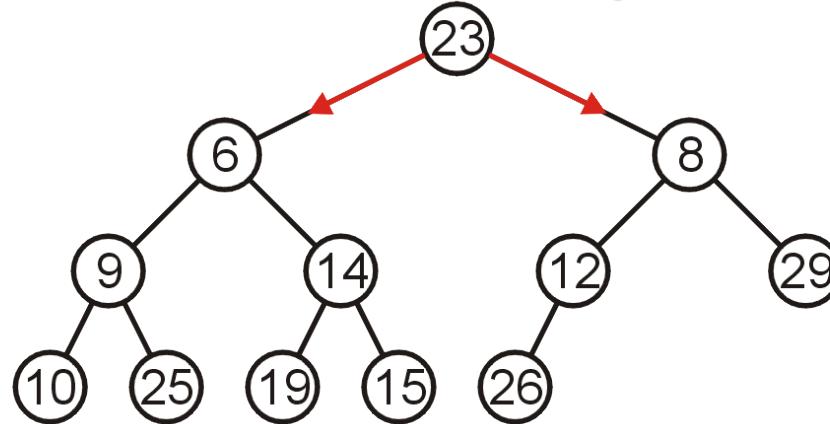
# Array Implementation: Pop

---

Now percolate down

Compare Node 1 with its children: Nodes 2 and 3

- Swap 23 and 6



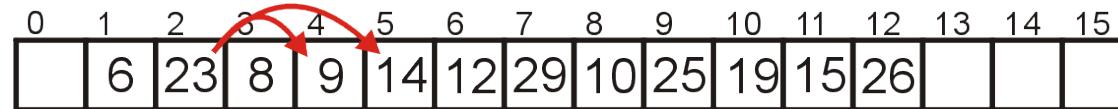
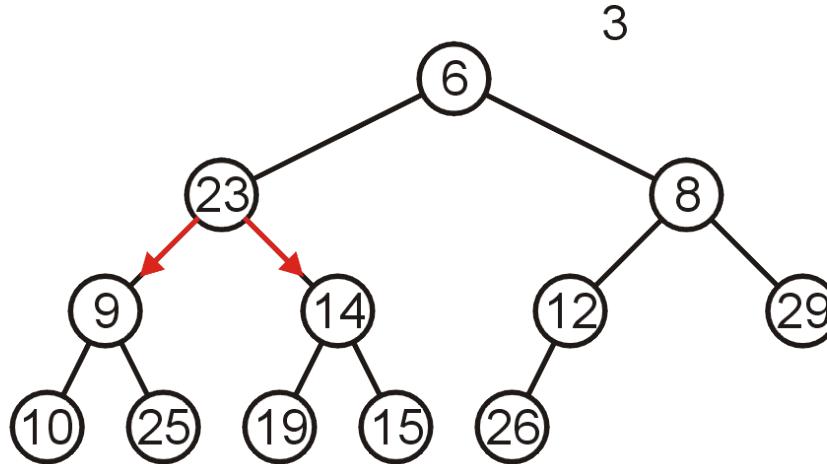
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	23	6	8	9	14	12	29	10	25	19	15	26			

# Array Implementation: Pop

---

Compare Node 2 with its children: Nodes 4 and 5

- Swap 23 and 9

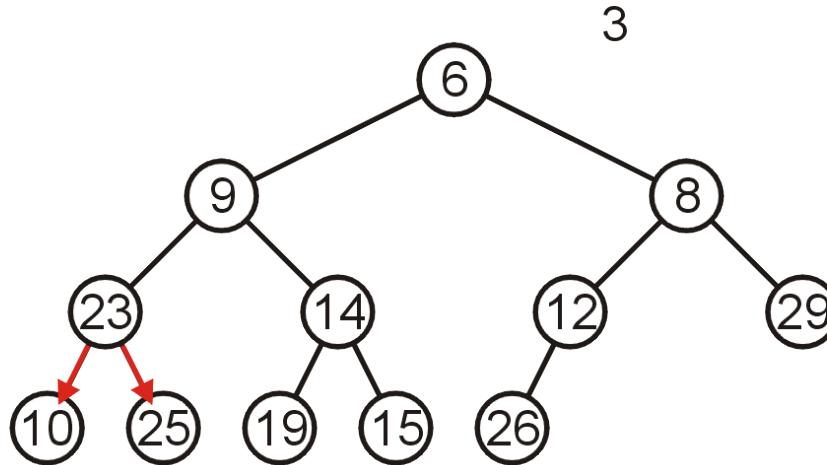


# Array Implementation: Pop

---

Compare Node 4 with its children: Nodes 8 and 9

- Swap 23 and 10

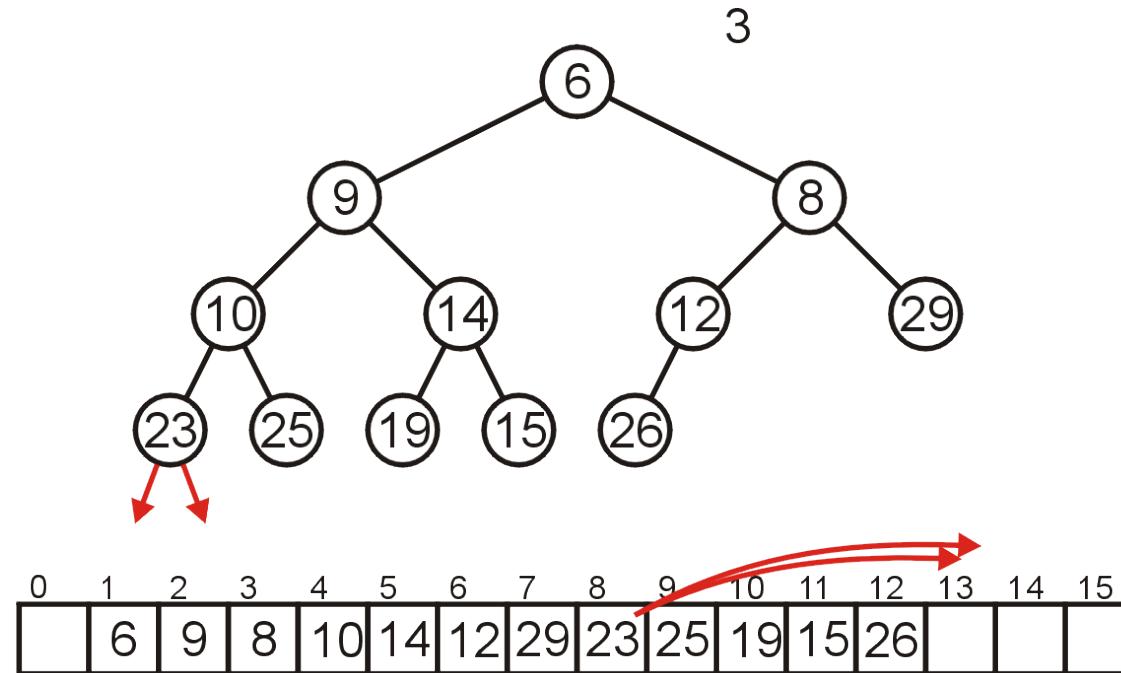


# Array Implementation: Pop

---

The children of Node 8 are beyond the end of the array:

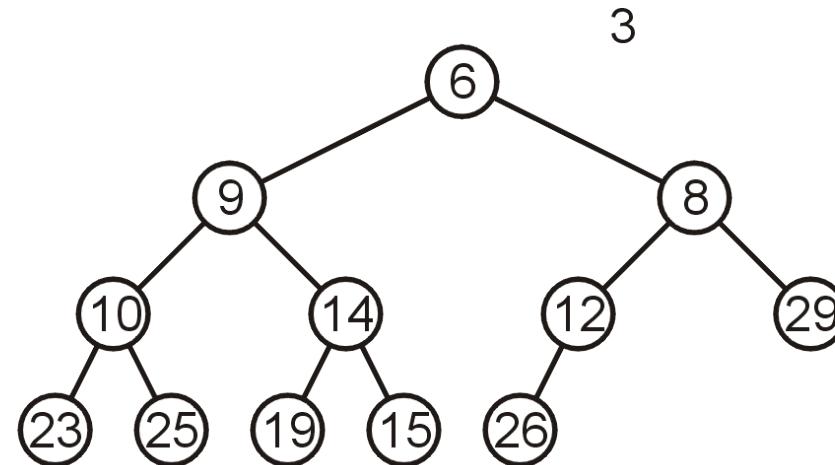
- Stop



# Array Implementation: Pop

---

The result is a binary min-heap



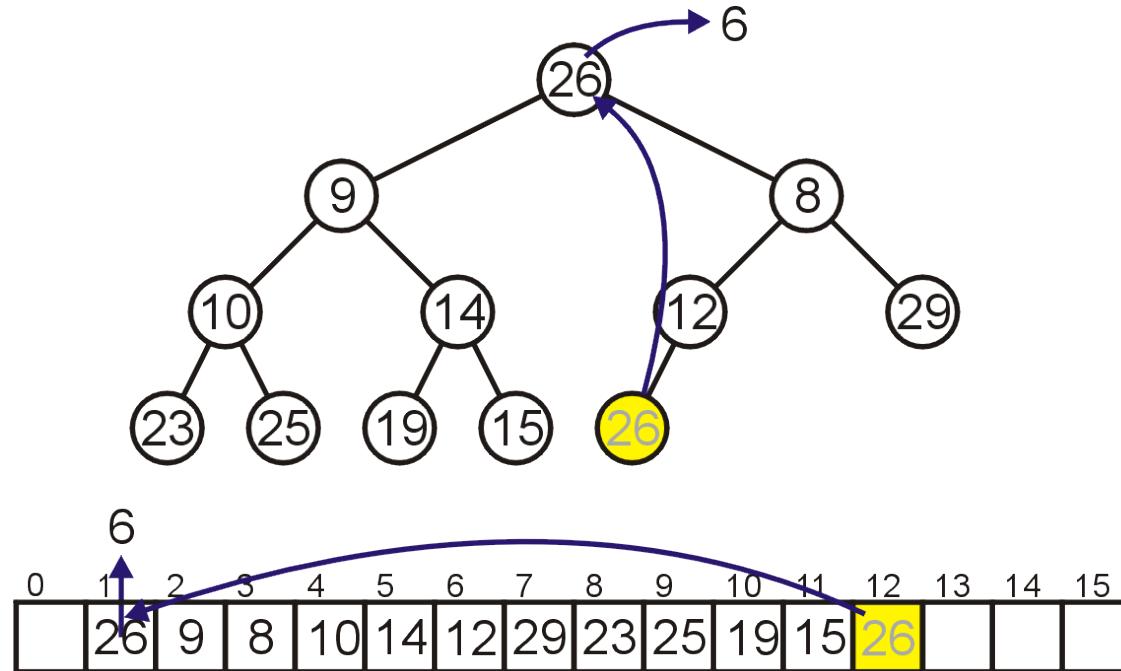
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	6	9	8	10	14	12	29	23	25	19	15	26			

# Array Implementation: Pop

---

Dequeuing the minimum again:

- Copy 26 to the root

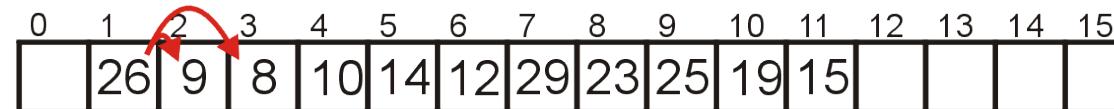
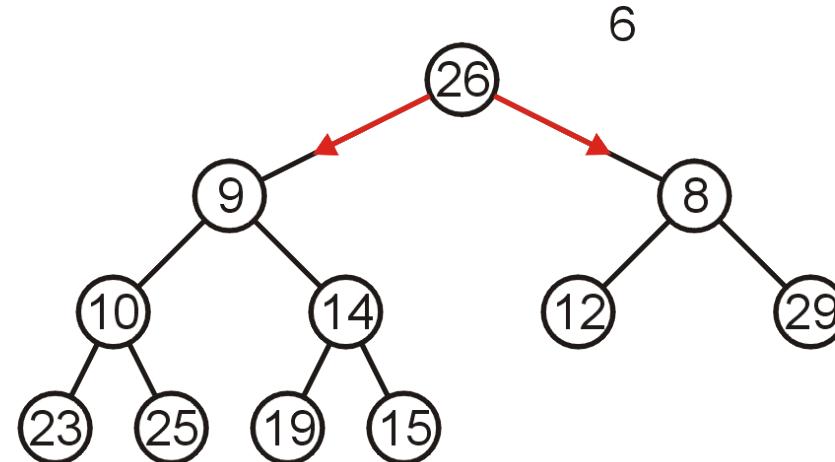


# Array Implementation: Pop

---

Compare Node 1 with its children: Nodes 2 and 3

- Swap 26 and 8

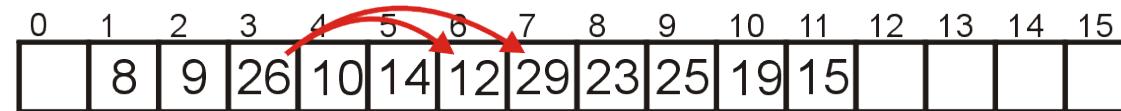
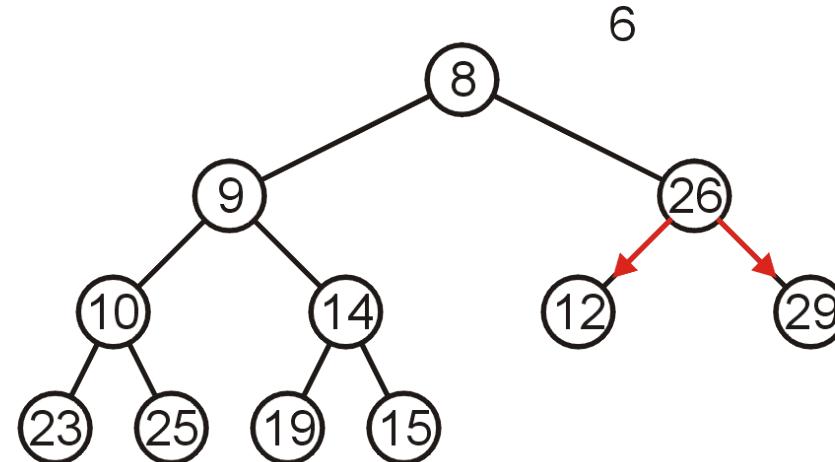


# Array Implementation: Pop

---

Compare Node 3 with its children: Nodes 6 and 7

- Swap 26 and 12

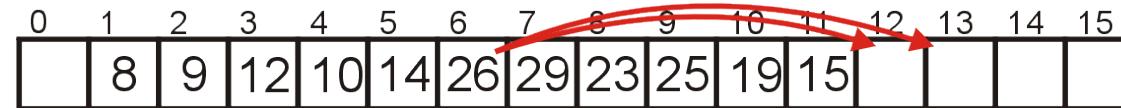
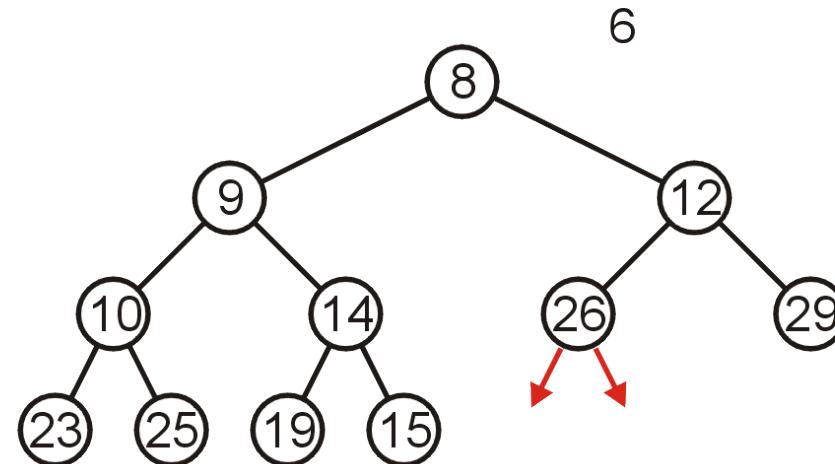


# Array Implementation: Pop

---

The children of Node 6, Nodes 12 and 13 are unoccupied

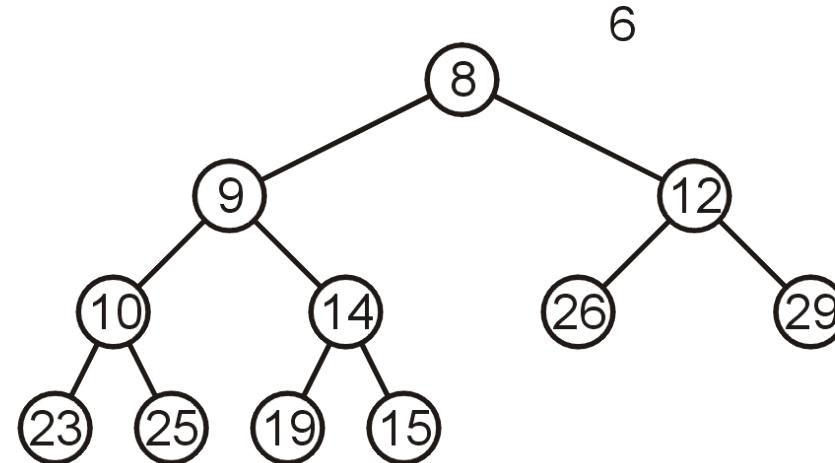
- Currently, count == 11



# Array Implementation: Pop

---

The result is a min-heap



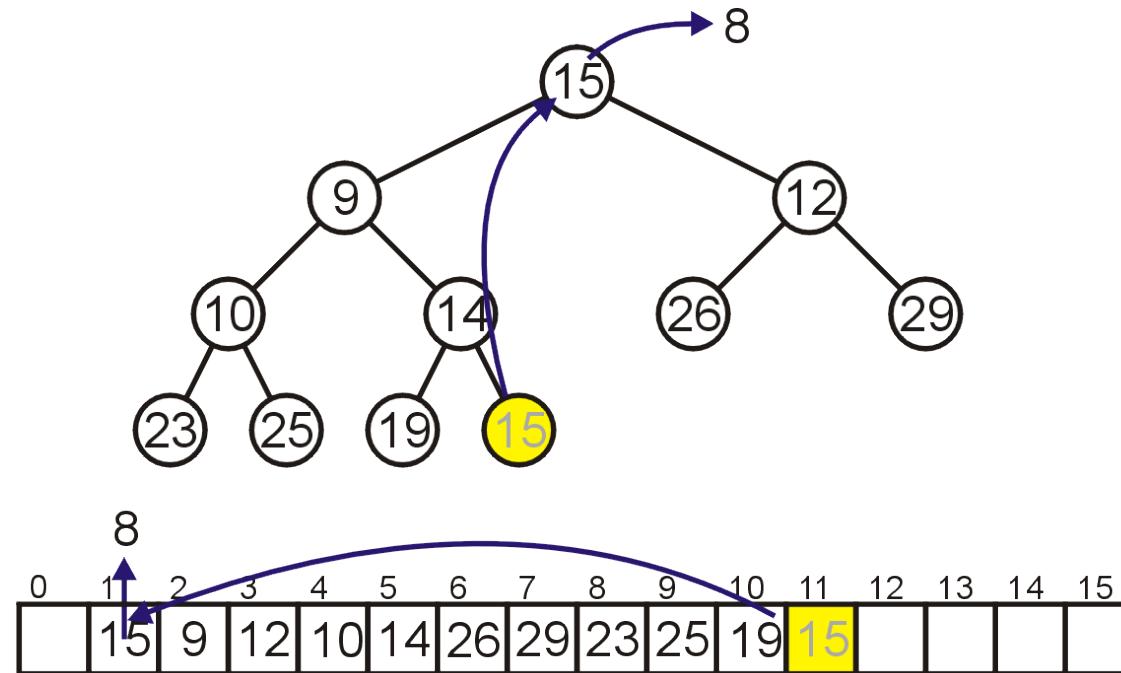
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	8	9	12	10	14	26	29	23	25	19	15				

# Array Implementation: Pop

---

Dequeuing the minimum a third time:

- Copy 15 to the root

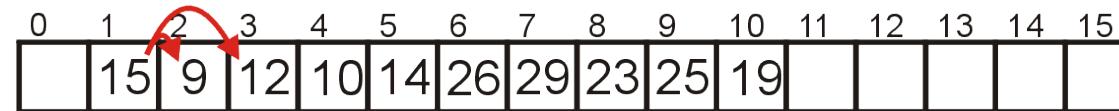
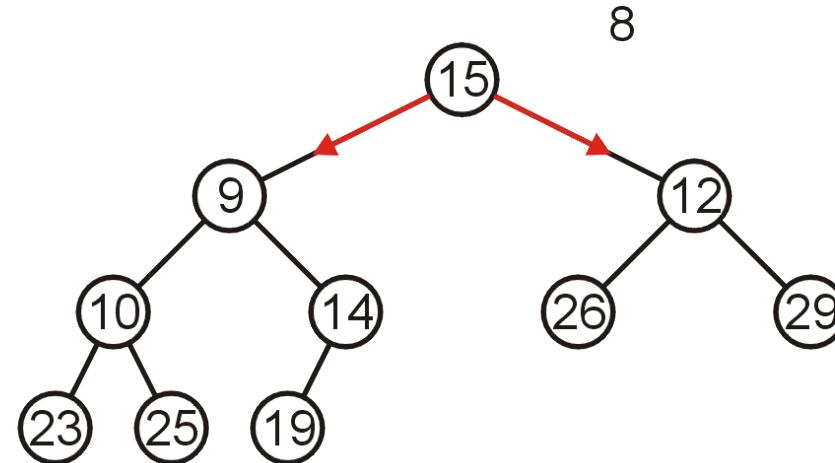


# Array Implementation: Pop

---

Compare Node 1 with its children: Nodes 2 and 3

- Swap 15 and 9

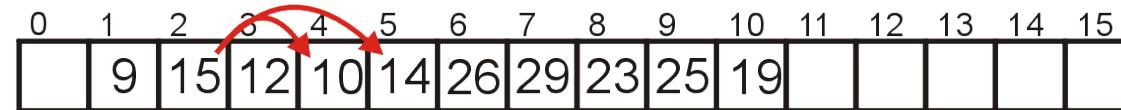
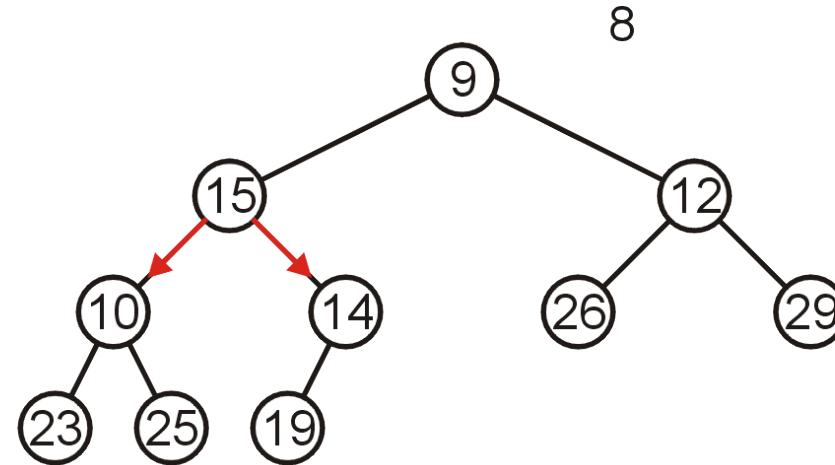


# Array Implementation: Pop

---

Compare Node 2 with its children: Nodes 4 and 5

- Swap 15 and 10

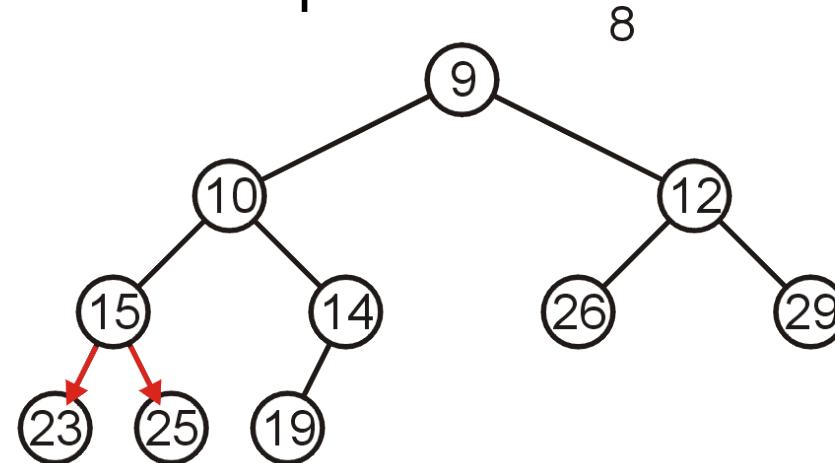


# Array Implementation: Pop

---

Compare Node 4 with its children: Nodes 8 and 9

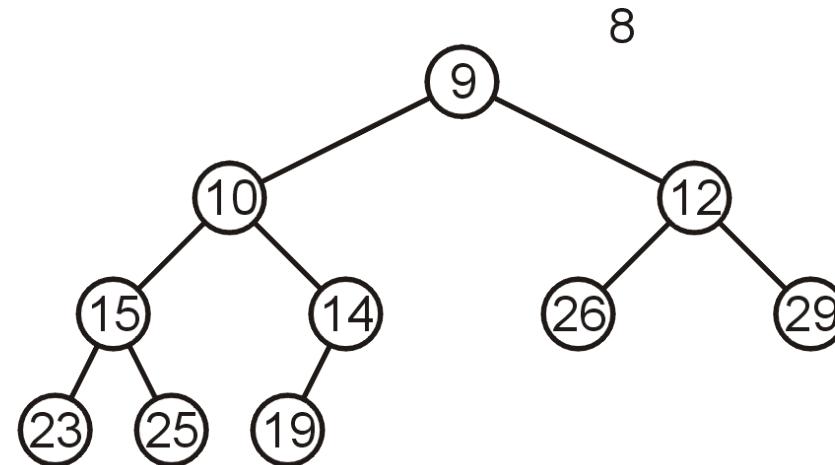
- $15 < 23$  and  $15 < 25$  so stop



# Array Implementation: Pop

---

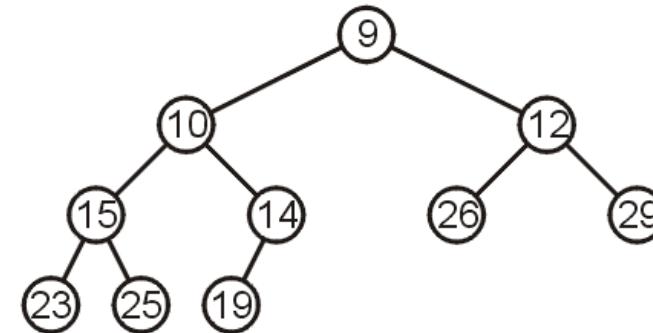
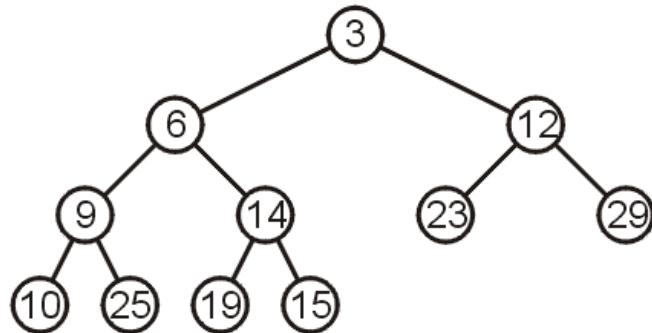
The result is a properly formed binary min-heap



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	9	10	12	15	14	26	29	23	25	19					

# Array Implementation: Pop

After all our modifications, the final heap is



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	3	6	12	9	14	23	29	10	25	19	15				

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	9	10	12	15	14	26	29	23	25	19					

# Run-time Analysis

---

Accessing the top object is  $\Theta(1)$

Popping the top object is  $O(\ln(n))$

- We copy something that is already in the lowest depth—it will likely be moved back to the lowest depth

How about push?

# Run-time Analysis

---

If we are inserting an object less than the root (at the front), then the run time will be  $\Theta(\ln(n))$

If we insert at the back (greater than any object) then the run time will be  $\Theta(1)$

How about an arbitrary insertion?

- It will be  $O(\ln(n))$ ? Could the average be less?

# Run-time Analysis

---

With each percolation, it will move an object past half of the remaining entries in the tree

- Therefore, after one percolation, it will probably be past half of the entries, and therefore *on average* will require no more percolations

$$\begin{aligned}\frac{1}{n} \sum_{k=0}^h (h-k)2^k &= \frac{2^{h+1} - h - 2}{n} \\ &= \frac{n-h-1}{n} = \Theta(1)\end{aligned}$$

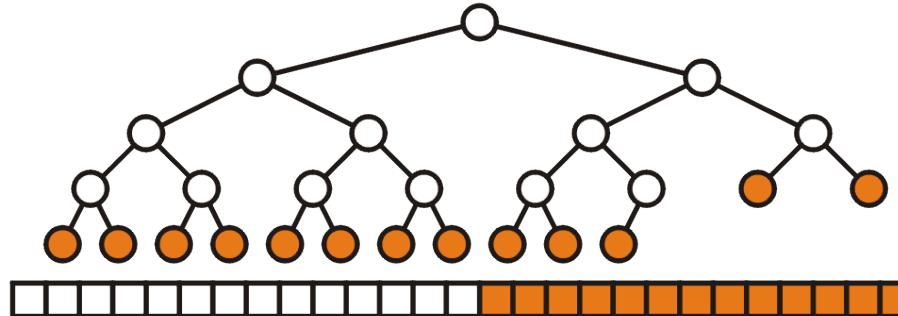
Therefore, we have an average run time of  $\Theta(1)$

# Run-time Analysis

---

An arbitrary removal requires that all entries in the heap be checked:  $O(n)$

A removal of the largest object in the heap still requires all leaf nodes to be checked – there are approximately  $n/2$  leaf nodes:  $O(n)$



# Run-time Analysis

---

Thus, our grid of run times is given by:

	front	arbitrary	back
insert	$O(\ln(n))^{*}$	$O(1)$	$O(1)$
access	$O(1)$	$O(n)$	$O(n)$
delete	$O(\ln(n))$	$O(n)$	$O(n)$

# Run-time Analysis

---

Some observations:

- Continuously inserting at the front of the heap (*i.e.*, the new object being pushed is less than everything in the heap) causes the run-time to drop to  $O(\ln(n))$
- If the objects are coming in order of priority, use a regular queue with swapping
- Merging two binary heaps of size  $n$  is a  $\Theta(n)$  operation

# Run-time Analysis

---

Other heaps have better run-time characteristics

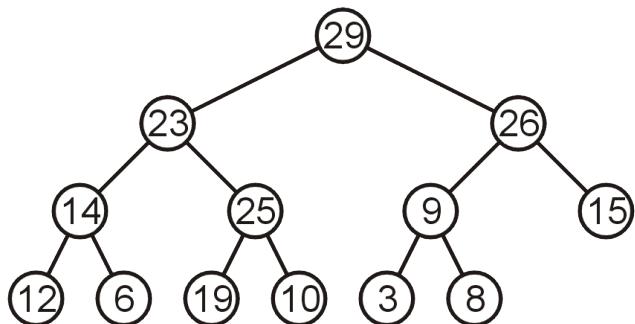
- Leftist, skew, binomial and Fibonacci heaps all use a node-based implementation requiring  $\Theta(n)$  additional memory
- For Fibonacci heaps, the run-time of all operations (including merging two Fibonacci heaps) except pop are  $\Theta(1)$

# Binary Max Heaps

---

A binary max-heap is identical to a binary min-heap except that the parent is always larger than either of the children

For example, the same data as before stored as a max-heap yields

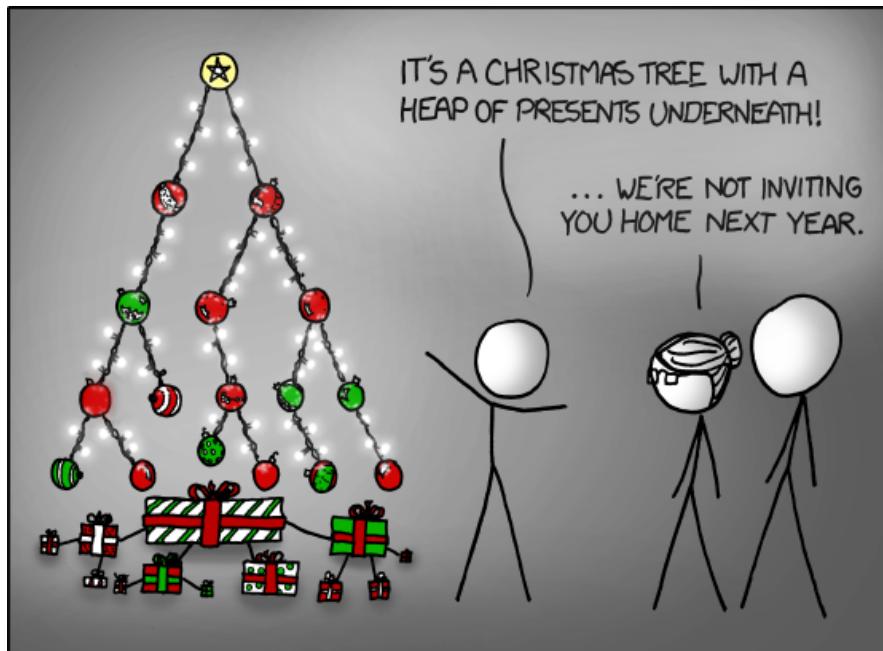


0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	29	23	26	14	25	9	15	12	6	19	10	3	8		

# Example

---

Here we have a max-heap of presents under a red-green tree:



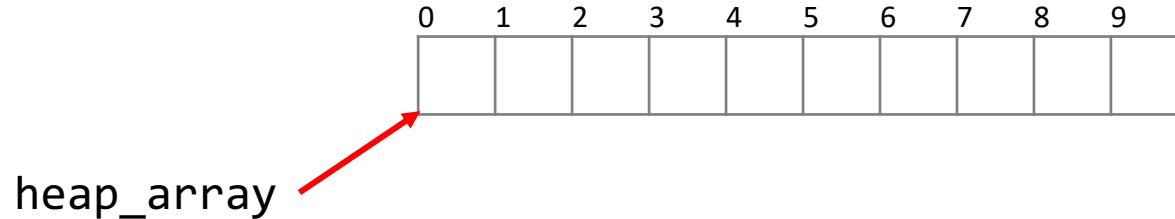
<http://xkcd.com/835/>

# Memory allocation and pointer arithmetic

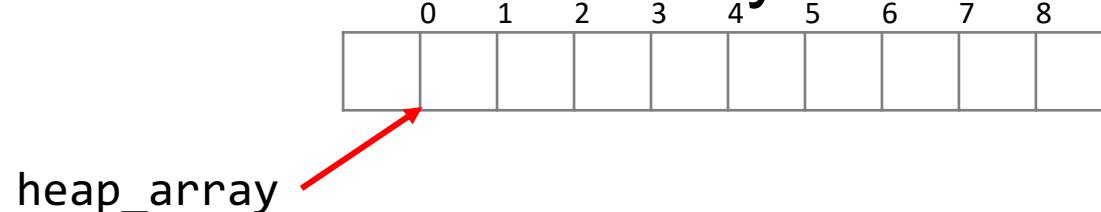
---

Do we really have to allocate one additional memory location for a binary tree-as-heap?

```
Type *heap_array = new Type[capacity() + 1];
```



Could we not just allocate one less memory and point to the previous location in memory

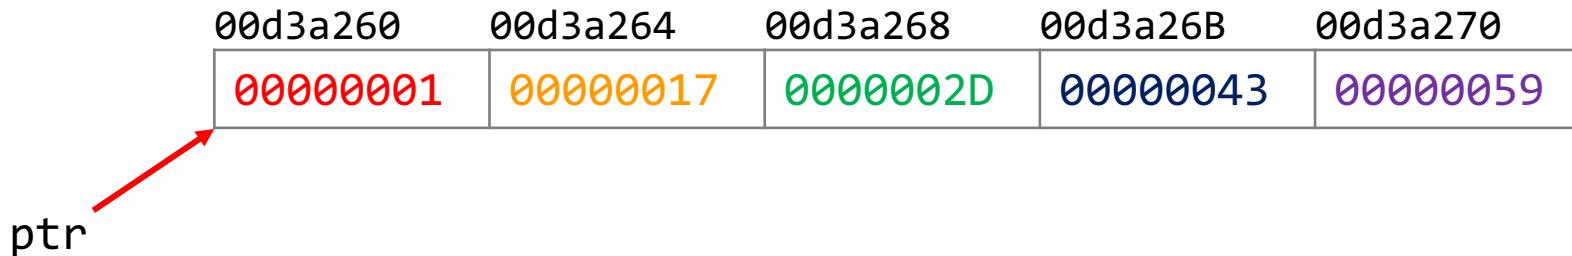


# Memory allocation and pointer arithmetic

---

To do this, we must understand pointer arithmetic:

```
int *ptr = new int[5] = {1, 23, 45, 67, 89};  
std::cout << ptr << std::endl;  
std::cout << *ptr << std::endl;
```



What is the output of?

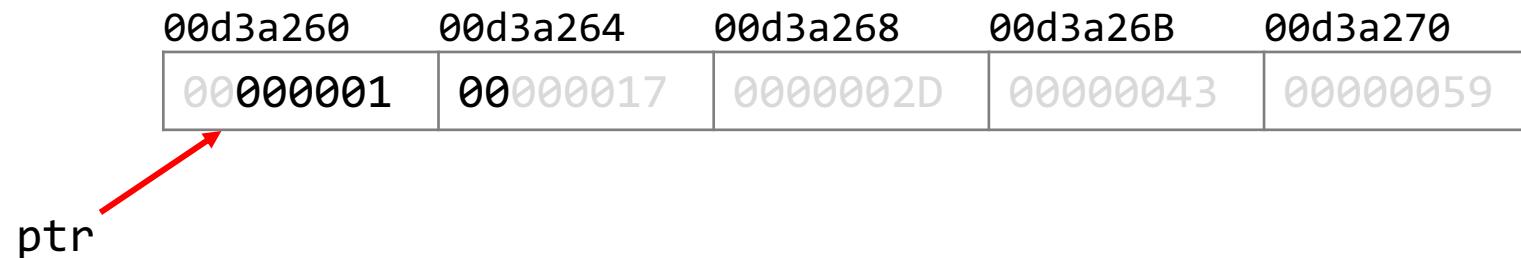
```
std::cout << (ptr + 1) << std::endl;  
std::cout << *(ptr + 1) << std::endl;
```

# Memory allocation and pointer arithmetic

---

Just adding one to the address would be, in almost all cases, useless

- Assuming big endian, this would have a value 256
- If this was little endian, it would be even more bizarre...



What is the output of?

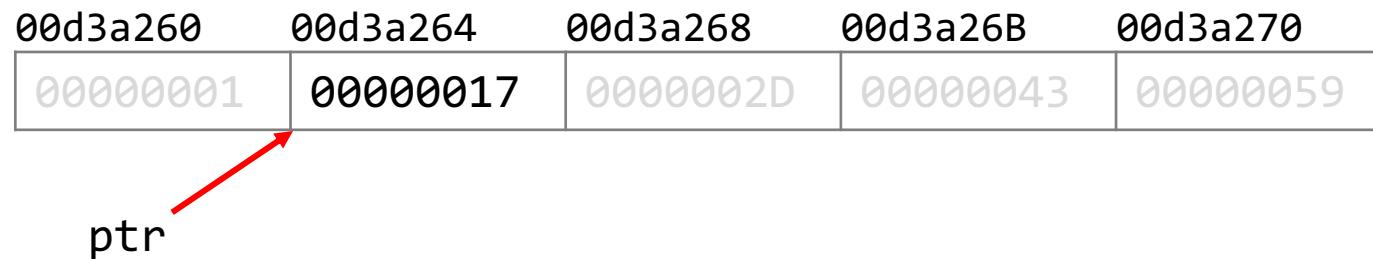
```
std::cout << (ptr + 1) << std::endl;  
std::cout << *(ptr + 1) << std::endl;
```

# Memory allocation and pointer arithmetic

---

Instead, C and C++ add as many bytes as the size of the object being pointed to

- In the cases of int, `sizeof( int ) == 4` on most 32-bit machines
- The output is 23



What is the output of?

```
std::cout << (ptr + 1) << std::endl;
std::cout << *(ptr + 1) << std::endl;
```

# Memory allocation and pointer arithmetic

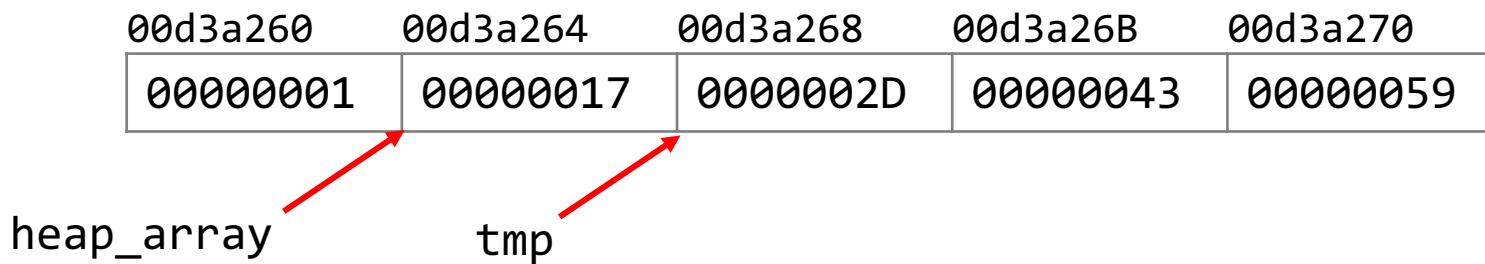
---

Essentially, these two statements are identical:

```
std::cout << ptr[i] << std::endl;  
std::cout << *(ptr + i) << std::endl;
```

Now you can do the following:

```
Type *tmp = new Type[capacity()];  
Type *heap_array = tmp - 1;
```



Now, `heap_array[1]`; and `tmp[0]`; both point to the same memory location

# Memory allocation and pointer arithmetic

---

## Issues:

- Never access or modify the contents of `heap_array[0]`
- When you deallocate memory, you must point to the original address returned by `new`:

```
delete [] (heap_array + 1);
```

- Pointer arithmetic is not for the faint of heart but it is fun; for example:

```
int array[N];
int *ptr = array;

// Print all the entries of the array
for ( int i = 0; i < N; ++i ) {
    std::cout << *(ptr++) << std::endl;
}
```

# Priority Queues

---

Now, does using a heap ensure that that object in the heap which:

- has the highest priority, and
- of that highest priority, has been in the heap the longest

Consider inserting seven objects, all of the same priority (colour indicates order):

2, 2, 2, 2, 2, 2, 2

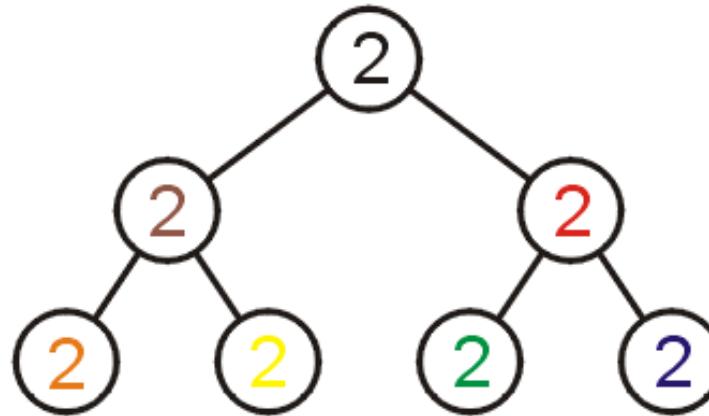
# Priority Queues

---

Whatever algorithm we use for promoting must ensure that the first object remains in the root position

- Thus, we must use an insertion technique where we only percolate up if the priority is lower

The result:



Challenge:

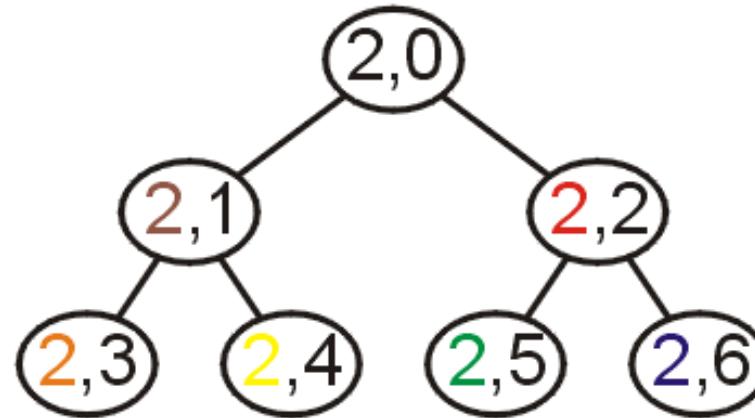
- Come up with an algorithm which removes all seven objects in the original order

# Lexicographical Ordering

---

A better solution is to modify the priority:

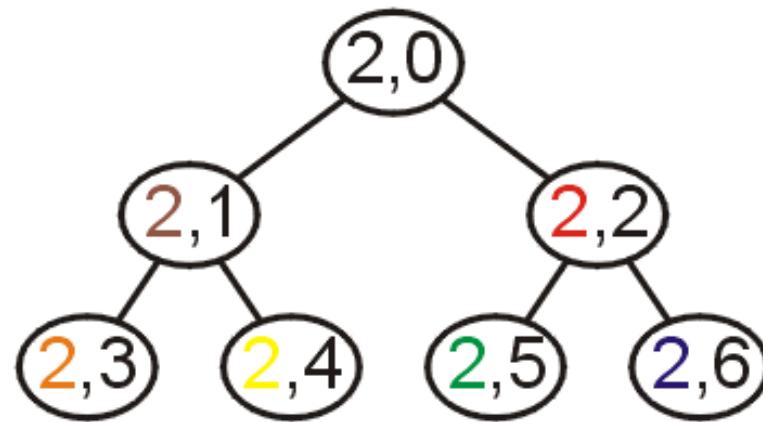
- Track the number of insertions with a counter  $k$  (initially 0)
- For each insertion with priority  $n$ , create a hybrid priority  $(n, k)$  where:  
 $(n_1, k_1) < (n_2, k_2)$  if  $n_1 < n_2$  or  $(n_1 = n_2 \text{ and } k_1 < k_2)$



# Priority Queues

---

Removing the objects would be in the following order:

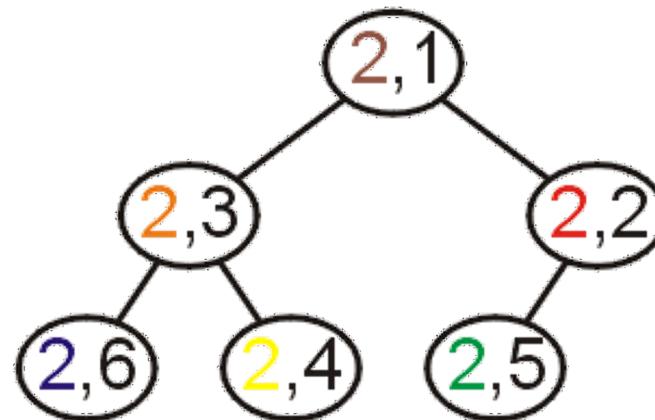


# Priority Queues

---

Popped: 2

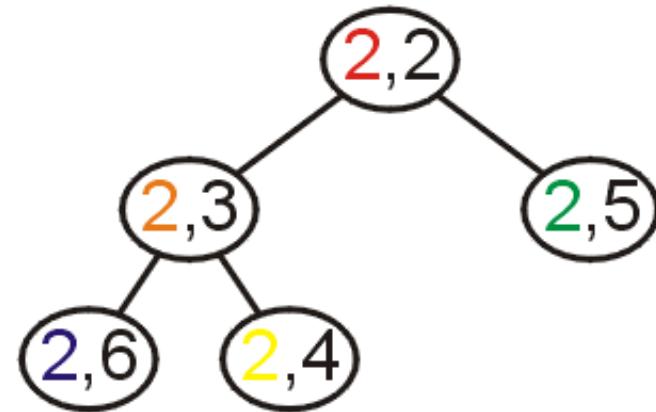
- First,  $(2,1) < (2, 2)$  and  $(2, 3) < (2, 4)$



# Priority Queues

---

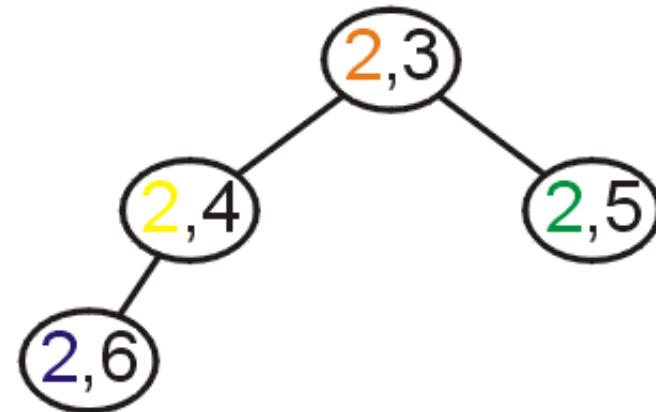
Removing the objects would be in the following order:



# Priority Queues

---

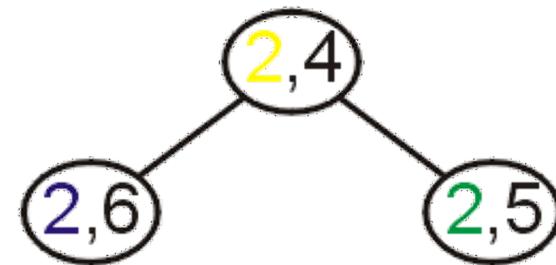
Removing the objects would be in the following order:



# Priority Queues

---

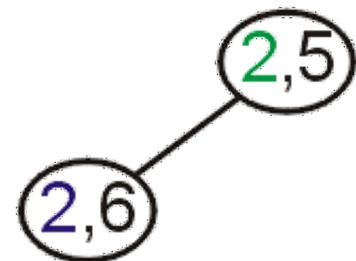
Removing the objects would be in the following order:



# Priority Queues

---

Removing the objects would be in the following order:



# Summary

---

In this talk, we have:

- Discussed binary heaps
- Looked at an implementation using arrays
- Analyzed the run time:
  - Head  $\Theta(1)$
  - Push  $\Theta(1)$  average
  - Pop  $O(\ln(n))$
- Discussed implementing priority queues using binary heaps
- The use of a lexicographical ordering